

# Algorithmics

## Parallel algorithms

Vicente García Díaz – [garciavicente@uniovi.es](mailto:garciavicente@uniovi.es)

*University of Oviedo, 2016*

# Table of contents

## *Parallel algorithms*

1. Basic concepts

2. Fork/Join Framework

- Examples



# *Basic concepts*

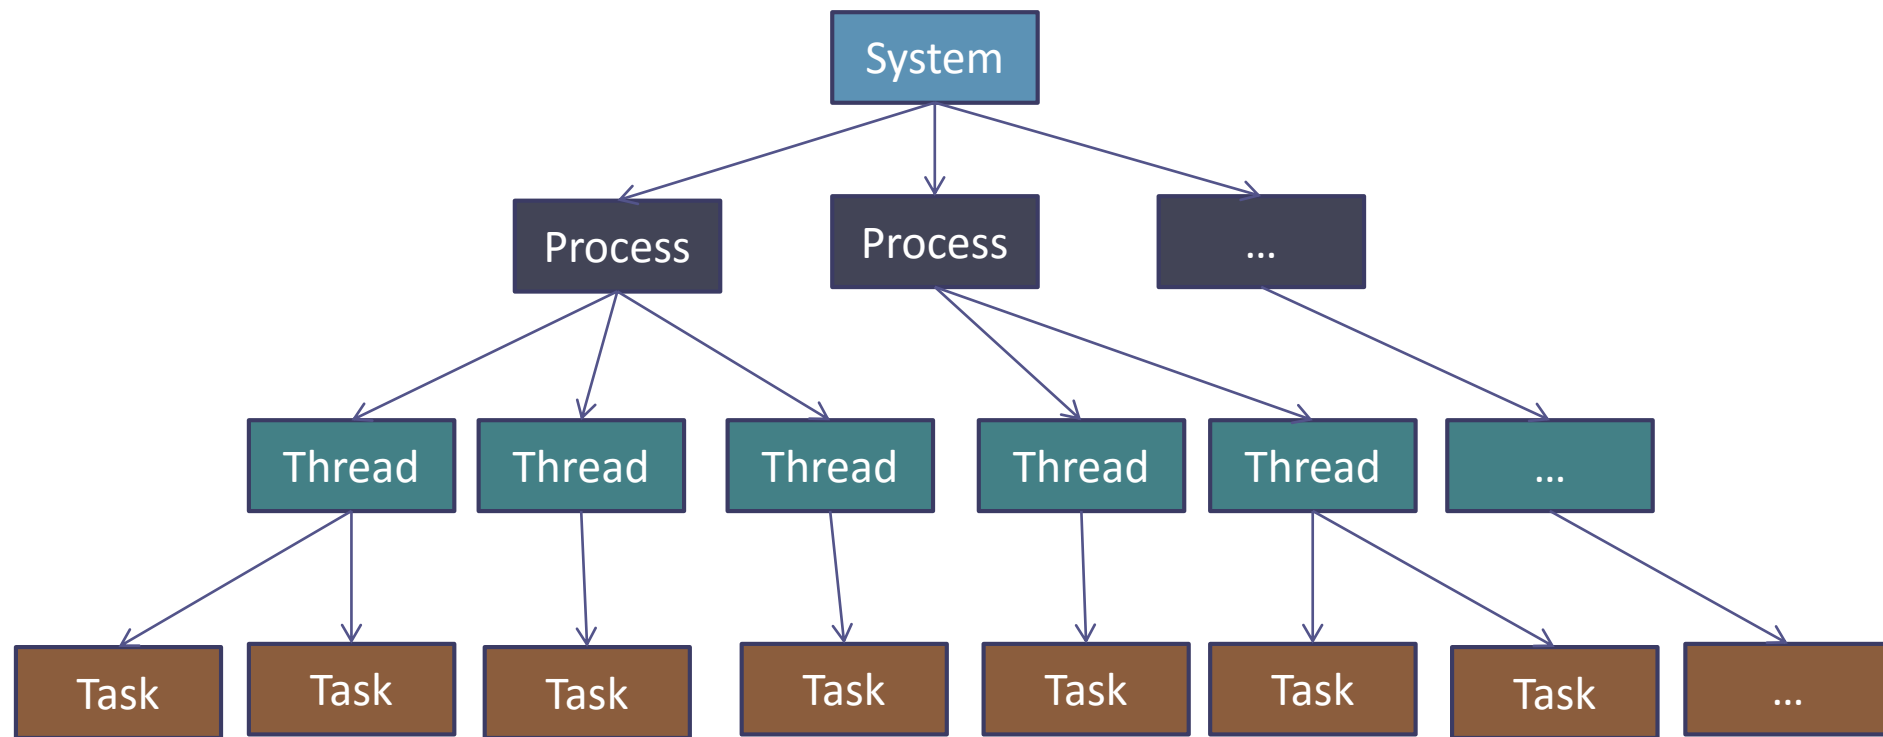
# Concurrent computing

- The concurrent computation focuses on performing two or more tasks on one or more CPUs
- It is the opposite to sequential computation
- Very common to perform tasks "simultaneously"
  - When we are waiting for the user to enter a value, we may be saving others in a database
  - Another example are the tasks that are performed in the background
- Actually they do not have to do the tasks at the same time because they may be using only one CPU
  - It gives the feeling of parallelism but sentences are executed one after the other in sequence
- There are used Java objects of type `Thread` to work with threads


# Mechanism to create concurrence (I)

- **Process**
  - An operating-system abstraction that allows one computer system to support different units of execution
- **Thread**
  - Belongs to exactly one process and is a block that is executed independently of others. Threads can share resources in the process
- **Task**
  - A component or small unit of execution
- The operating system is responsible for executing the threads concurrently on one or more processors (concurrently or in parallel)
- You have to synchronize the threads so that threads do not block each other and do not try to simultaneously access certain system resources

# Mechanism to create concurrency (II)



# Parallel computing



- It is intrinsically concurrent (multithreaded)
- Used to perform simultaneous tasks using more than one CPU
- Can greatly speed up various operations
  - P.e., sorting, searching, transformations, ...
- It is definitely true for any task that can be broken down into smaller components (divide and conquer)

The background of the slide is a light-colored marbled paper with a repeating pattern of small, stylized floral or leaf motifs. At the top, there is a solid teal header bar. On the right side of this bar, there are two horizontal white lines of different lengths, one above the other.

# *Fork/Join Framework*



# Concept

- It is available since Java 7
  - `java.util.concurrent`
- The aim is to utilize the full power of parallel processing to improve application performance
- General idea:

```
if (work is small enough)
    Do the work directly
else{
    Break up the work into smaller components
    Perform the resolution of components in parallel
    Wait for results
}
```

# Advantages

1. It facilitates the creation and use of multiple threads
2. It makes use of multiple processors in parallel automatically

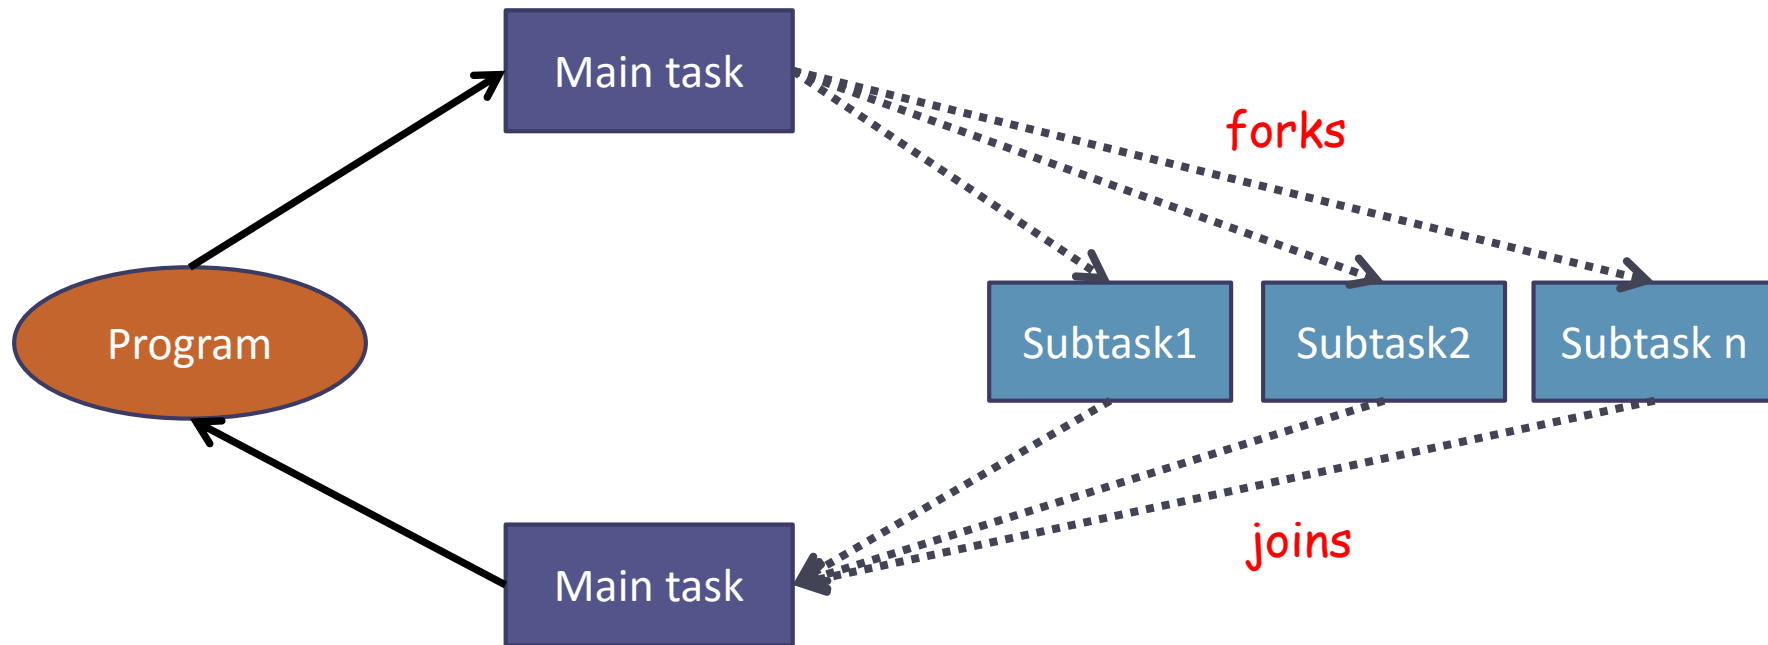
# Pseudocode

```
public <T> solve(Problem problem){
    if (problem.size < SEQUENTIAL_THRESHOLD)
        return solveSequentially(problem);
    else {
        <T> left, right;
        INVOKE-IN-PARALLEL {
            left = solve(extractLeftHalf(problem));
            right = solve(extractRightHalf(problem));
        }
        return combine(left, right);
    }
}
```

# Approach (I)

1. Partition into subproblems
  - Divide a large problem into smaller ones
2. Create subtasks
  - Design a solution for each of the subproblems independently (through a thread)
3. Fork subtasks
  - Indicate that you want to start solving subproblems (the threads are sent to a pool of threads)
  - The pool size depends on the number of CPUs and other considerations
4. Join subtasks
  - Wait for processing and solving each of the subproblems (usually all)
5. Compose solution
  - Compose the solution from the obtained partial solutions

# Approach (II)



# Main components

- `ForkJoinTask<V>`
  - An abstract class that defines a task
- `ForkJoinPool`
  - A class that manages the execution of `ForkJoinTasks`
- `RecursiveAction`
  - A subclass of `ForkJoinTask<V>` for tasks that do not return values
- `RecursiveTask<V>`
  - A subclass of `ForkJoinTask<V>` for tasks that return values

## ForkJoinTask<V> (I)

- It is an abstract class that defines a class that can be managed by `ForkJoinPool`
- `V` specifies the result type of the task
- `ForkJoinTask` raises the level of abstraction of `Thread`
  - Task VS Thread of execution
- `ForkJoinTasks` are executed by threads managed by a thread pool of type `ForkJoinPool`
- `ForkJoinTasks` are much more efficient than threads

## ForkJoinTask<V> (II)

- `final ForkJoinTask<V> fork()`
  - Submits the invoking task for asynchronous execution
- `final V join()`
  - Waits until the task on which it is called terminates
- `final V invoke()`
  - Combines the `fork` and `join` operations into a single call
- `static void invokeAll(ForkJoinTask<?> task1, ForkJoinTask<?> task2)`
- `static void invokeAll(ForkJoinTask<?>... taskList)`



# RecursiveAction

- It is a subclass of `ForkJoinTask`
- Encapsulates a task that does not return values (`void`)
- `protected abstract void compute()`
- It is used to implement a recursive, **divide-and-conquer** strategy

# RecursiveTask<V>

- It is a subclass of ForkJoinTask
- Encapsulates a task that returns a result
- `protected abstract V compute()`
- It is used to implement a recursive, **divide-and-conquer** strategy

# ForkJoinPool (I)

- It is responsible for launching and managing the execution of `ForkJoinTasks`
- Two very typical constructors
  - `ForkJoinPool()`
    - Supports a level of parallelism equal to the number of processors available in the system
  - `ForkJoinPool(int pLevel)`
    - Lets you specify the level of parallelism
- Level of parallelism
  - Number of threads that can be executed concurrently / in parallel

## ForkJoinPool (II)

- The level of parallelism does not limit the number of tasks that can be managed by the pool
- ...although the number of tasks that can execute simultaneously cannot exceed the number of processors
- ...and the level of parallelism is a target, not a guarantee
- `<T> invoke(ForkJoinTask<T> task)`
  - The calling code waits until the return of the method
- `void execute(ForkJoinTask<?> task)`
  - The calling code continues execution asynchronously

# Divide and Conquer strategy

- The idea is to make recursive calls to classes that extend `RecursiveTask` or `RecursiveAction`
- Reducing the size of the problem to a size that can be managed sequentially faster than by creating another division
- The key is to find a good **threshold** to stop the decomposing of the problem
- We must find a balance between what it costs to break the problem and what it costs to solve it
- It **cannot** be based on the **number of processors** because it changes in different computers and other applications may be making use of them

# Problem

76	13	97	90	48	76	94	59	42	5
----	----	----	----	----	----	----	----	----	---



5776	169	9409	8100	2304	5776	8836	3481	1764	25
------	-----	------	------	------	------	------	------	------	----

## Fork/Join Framework

RecursiveAction. Obtaining the square of the values of an array

```
package com.vgd.algorithmsparallel;

import java.util.concurrent.*;

public class RecursiveActionSquare extends RecursiveAction {
    private static final long serialVersionUID = 1L;

    int[] data; //Array with numbers (data)
    int start, end; //Determines what part of data to process.
    //In real word code, its optimal value can be determined by experimentation
    final int seqThreshold = 100; //Arbitrary set at 100

    RecursiveActionSquare(int[] data, int start, int end) {
        this.data = data;
        this.start = start;
        this.end = end;
    }

    @Override
    protected void compute() {
        //If number of elements is below the sequential threshold, then process sequentially
        if((end - start) < seqThreshold) {
            for(int i = start; i < end; i++) {
                data[i] = data[i]*data[i]; //Transform each element into its square
            }
        }
        else { //Continue to break the data into smaller components
            int middle = (start + end) / 2; //Find the midpoint
            //Invoke new subtasks
            invokeAll(new RecursiveActionSquare(data, start, middle),
                    new RecursiveActionSquare(data, middle, end));
        }
    }
}
```

## Fork/Join Framework

RecursiveAction. Obtaining the square of the values of an array

```
package com.vgd.algorithmsparallel;

import java.util.Random;
import java.util.concurrent.*;

public class RecursiveActionSquareTest {
    public static void main(String[] args) {
        ForkJoinPool pool = new ForkJoinPool(); //Task pool
        Random rnd = new Random(); //Random numbers
        int[] data = new int[1000]; //Numbers to work with

        for(int i = 0; i < data.length; i++) //Some values
            data[i] = rnd.nextInt(100);

        System.out.println("The original sequence:");
        for(int i=0; i < data.length; i++)
            System.out.print(data[i] + " ");
        System.out.println("\n");

        RecursiveActionSquare task = new RecursiveActionSquare(data, 0, data.length);
        pool.invoke(task); //Start the main ForkJoinTask

        System.out.println("The transformed sequence:");
        for(int i=0; i < data.length; i++)
            System.out.print(data[i] + " ");
        System.out.println();
    }
}
```





## Fork/Join Framework

RecursiveAction. Comparison of different thresholds and CPUs

```

public class RecursiveActionComparison extends RecursiveAction {
    private static final long serialVersionUID = 1L;

    int[] data; //Array with numbers (data)
    int start, end; //Determines what part of data to process.
    //In real word code, its optimal value can be determined by experimentation
    int threshold = 100; //Arbitrary set at 1000

    RecursiveActionComparison(int[] data, int start, int end, int threshold) {
        this.data = data;
        this.start = start;
        this.end = end;
        this.threshold = threshold;
    }

    @Override
    protected void compute() {
        //If number of elements is below the sequential threshold, then process sequentially
        if((end - start) < threshold) {
            //Time consuming task so the effects of concurrent execution are more observable
            for(int i = start; i < end; i++) {
                data[i] = (int)(Math.cbrt(data[i]));
            }
        }
        else { //Continue to break the data into smaller components
            int middle = (start + end) / 2; //Find the midpoint
            //Invoke new subtasks
            invokeAll(new RecursiveActionComparison(data, start, middle, threshold),
                      new RecursiveActionComparison(data, middle, end, threshold));
        }
    }
}

```

## Fork/Join Framework

RecursiveAction. Comparison of different thresholds and CPUs

```
public class RecursiveActionComparisonTest {
    public static void main(String[] args) {
        int level; //Level of parallelism
        int threshold;

        if(args.length != 2) {
            System.out.println("Usage: RecursiveActionComparisonTest threshold parallism");
            return;
        }

        level = Integer.parseInt(args[0]);
        threshold = Integer.parseInt(args[1]);

        ForkJoinPool pool = new ForkJoinPool(level); //Task pool with the level
        Random rnd = new Random(); //Random numbers
        int[] data = new int[10000000]; //Numbers to work with

        for(int i = 0; i < data.length; i++) //Some values
            data[i] = rnd.nextInt(100);

        RecursiveActionComparison task =
            new RecursiveActionComparison(data, 0, data.length, threshold);

        long t1 = System.currentTimeMillis(); //to measure the time
        pool.invoke(task); //Start the main ForkJoinTask
        long t2 = System.currentTimeMillis();

        System.out.println("Level of parallelism: " + level);
        System.out.println("Sequential threshold: " + threshold);
        System.out.println("Elapsed time: " + (t2-t1) + " ms");
        System.out.println();
    }
}
```

# Get information about the parallelism

- You can know the current level of parallelism and the number of processors available on the computer by calling two simple methods:

```
public class ParallelismInfo {  
    public static void main(String[] args) {  
        ForkJoinPool pool = new ForkJoinPool(5); //Task pool  
  
        System.out.println("Level of parallelism: " +  
            pool.getParallelism());  
  
        System.out.println("Available processors: " +  
            Runtime.getRuntime().availableProcessors());  
    }  
}
```

# Problem

76	13	97	90	48	76	94	59	42	52
----	----	----	----	----	----	----	----	----	----



647

## Fork/Join Framework

RecursiveTask<V>. Sum the elements of an array

```
public class RecursiveTaskSum extends RecursiveTask<Double> {
    private static final long serialVersionUID = 1L;

    double[] data; //Array with numbers (data)
    int start, end; //Determines what part of data to process.
    int threshold = 10000; //Arbitrary set

    RecursiveTaskSum(double[] data, int start, int end) {
        this.data = data;
        this.start = start;
        this.end = end;
    }

    @Override
    protected Double compute() {
        double sum = 0;
        if((end - start) < threshold) {
            for(int i = start; i < end; i++) {
                sum += data[i];
            }
        }
        else { //Continue to break the data into smaller components
            int middle = (start + end) / 2; //Find the midpoint
            //Invoke new subtasks
            RecursiveTaskSum subTaskA = new RecursiveTaskSum(data, start, middle);
            RecursiveTaskSum subTaskB = new RecursiveTaskSum(data, middle, end);
            subTaskA.fork(); //Start each subtask by forking
            //Wait for the subtasks to return, and aggregate the results
            sum = subTaskB.compute() + subTaskA.join();
            //sum = subTaskA.invoke() + subTaskB.invoke();
        }
        return sum;
    }
}
```

```
public class RecursiveTaskSumTest {
    public static void main(String[] args) {
        ForkJoinPool pool = new ForkJoinPool(); //Task pool with the level
        double[] data = new double[9999999]; //Numbers to work with


        //Initialize nums with values that alternate between positive and negative.
        for(int i=0; i < data.length; i++)
            data[i] = (double)(((i%2) == 0) ? i : -i) ;


        RecursiveTaskSum task =
            new RecursiveTaskSum(data, 0, data.length);


        long t1 = System.currentTimeMillis(); //to measure the time
        double result = pool.invoke(task); //Start the main ForkJoinTask
        long t2 = System.currentTimeMillis();


        System.out.println("Elapsed time: " + (t2-t1) + " ms");
        System.out.println("Result: " + result);
    }
}
```


# Alternatives for obtaining task results

 Temp subTaskA = new Temp(data, start, middle);  
 Temp subTaskB = new Temp(data, middle, end);  
 subTaskA.fork();  
 subTaskB.fork();  
 sum = subTaskA.join() + subTaskB.join();

 Temp subTaskA = new Temp(data, start, middle);  
 Temp subTaskB = new Temp(data, middle, end);  
 sum = subTaskA.invoke() + subTaskB.invoke();

 Temp subTaskA = new Temp(data, start, middle);  
 Temp subTaskB = new Temp(data, middle, end);  
 subTaskA.fork();  
 sum = subTaskA.join() + subTaskB.compute();

 Temp subTaskA = new Temp(data, start, middle);  
 Temp subTaskB = new Temp(data, middle, end);  
 subTaskA.fork();  
 sum = subTaskA.join() + subTaskB.invoke();

 Temp subTaskA = new Temp(data, start, middle);  
 Temp subTaskB = new Temp(data, middle, end);  
 subTaskA.fork();  
 subTaskB.fork();  
 sum = subTaskB.join() + subTaskA.join();

Temp subTaskA = new Temp(data, start, middle);  
 Temp subTaskB = new Temp(data, middle, end);  
 subTaskA.fork();  
 sum = subTaskB.compute() + subTaskA.join();


Temp subTaskA = new Temp(data, start, middle);  
 Temp subTaskB = new Temp(data, middle, end);  
 subTaskA.fork();  
 sum = subTaskB.invoke() + subTaskA.join();



# Other methods of ForkJoinTask

- To cancel a task
  - `boolean cancel(boolean interruptOK)`
- To know whether a task has been canceled
  - `boolean isCancelled()`
- To determine the state of completeness of a task
  - `final boolean isCompletedNormally()`
  - `final boolean isCompletedAbnormally()`
  - `final boolean isCompleted()`

# Problem



# Classic Fibonacci Algorithm

```

public class FibonacciAlgorithm {
    public int n;

    public FibonacciAlgorithm(int n){
        this.n = n;
    }

    public long solve(){
        return fibonacci(this.n);
    }

    private long fibonacci(int n) {
        if (n <= 1)
            return n;
        else
            return fibonacci(n-1) + fibonacci(n-2);
    }
}

public class FibonacciAlgorithmTest {

    public static void main(String[] args) throws Exception {
        int n = 10;
        FibonacciAlgorithm problem = new FibonacciAlgorithm(n);

        long t1 = System.currentTimeMillis();
        long result = problem.solve();
        long t2 = System.currentTimeMillis();

        System.out.println("Fibonacci problem: " + n);
        System.out.println("Result: " + result);
        System.out.println("Elapsed time: " + (t2-t1) + " ms");
    }
}

```

# Algorithm with RecursiveTask<V>

```

public class FibonacciTask extends RecursiveTask<Long> {
    private static final long serialVersionUID = 1L;
    private static final int THRESHOLD = 9;
    private FibonacciAlgorithm problem;
    public long result;

    public FibonacciTask(FibonacciAlgorithm problem) {
        this.problem = problem;
    }

    @Override
    public Long compute() {
        if (problem.n < THRESHOLD) {
            result = problem.solve();
        }
        else {
            FibonacciTask subTask1 = new
                FibonacciTask(new FibonacciAlgorithm(problem.n-1));
            FibonacciTask subTask2 = new
                FibonacciTask(new FibonacciAlgorithm(problem.n-2));
            subTask1.fork();
            result = subTask2.compute() + subTask1.join();
        }
        return result;
    }
}

public class FibonacciTaskTest {
    public static void main(String[] args) throws Exception {
        int n = 10;
        FibonacciAlgorithm problem = new FibonacciAlgorithm(n);
        FibonacciTask task = new FibonacciTask(problem);
        ForkJoinPool pool = new ForkJoinPool();

        long t1 = System.currentTimeMillis();
        long result = pool.invoke(task);
        long t2 = System.currentTimeMillis();

        System.out.println("Fibonacci problem: " + n);
        System.out.println("Result: " + result);
        System.out.println("Elapsed time: " + (t2-t1) + " ms");
    }
}

```

# Problem

- File processing tasks are favorable to be performed in parallel
- The number of files and the information contained in them can be very high
- For example:
  - Version management systems
  - Systems that process data
  - Searchers
  - ...



## Fork/Join Framework

## Processing files concurrently

```

class FileProcessingTask extends RecursiveAction {
    private static final long serialVersionUID = 1L;
    private static final int THRESHOLD = 5;
    List<File> javaFiles = null;
    String dirPath;

    public FileProcessingTask(String dirPath, List<File> javaFiles) {
        this.dirPath = dirPath;
        this.javaFiles = javaFiles;
    }


    @Override
    protected void compute() {
        if (javaFiles == null) { //First time to start processing files
            javaFiles = new ArrayList<File>();
            File sourceDir = new File(dirPath);
            if (sourceDir.isDirectory()) {
                for (File file : sourceDir.listFiles()){
                    javaFiles.add(file);
                }
            }
        }
        if (javaFiles.size() <= THRESHOLD) {
            processFiles(javaFiles);
        }
        else {
            int center = javaFiles.size() / 2;
            List<File> part1 = javaFiles.subList(0, center);
            List<File> part2 = javaFiles.subList(center, javaFiles.size());
            invokeAll(new FileProcessingTask(dirPath, part1),
                      new FileProcessingTask(dirPath, part2));
        }
    }

    protected void processFiles(List<File> filesToProcess) {
        for (File file : filesToProcess){
            System.out.println(Thread.currentThread().getName()
                               + " " + file.getName());
        }
    }
}

```


```
public class FileProcessingTasksTest {  
    public static void main(String[] args) {  
        FileProcessingTask problem = new FileProcessingTask("D:\\test\\source", null);  
        ForkJoinPool pool = new ForkJoinPool();  
  
        long t1 = System.currentTimeMillis();  
        pool.invoke(problem);  
        long t2 = System.currentTimeMillis();  
  
        System.out.println("Elapsed time: " + (t2-t1) + " ms");  
    }  
}
```

# Problem





# Fork/Join Framework

 Visualizing the Java Concurrent API - JDK 1.7.0\_05

Legend   Concurrency Training   Executors   Future   ReentrantLock   Condition   Semaphore   ReadWriteLock   CountDownLatch   CyclicBarrier   Phaser   Blinking Phaser   CompletionService   ConcurrentHashMap   Fork/Join   Help

Legend  
 Concurrency Training  
 Executors  
 Future  
 ReentrantLock  
 Condition  
 Semaphore  
 ReadWriteLock  
 CountDownLatch  
 CyclicBarrier  
 Phaser  
 Blinking Phaser  
 AtomicInteger  
 BlockingQueue  
 TransferQueue  
 CompletionService  
 ConcurrentHashMap  
 Fork/Join  
 Help

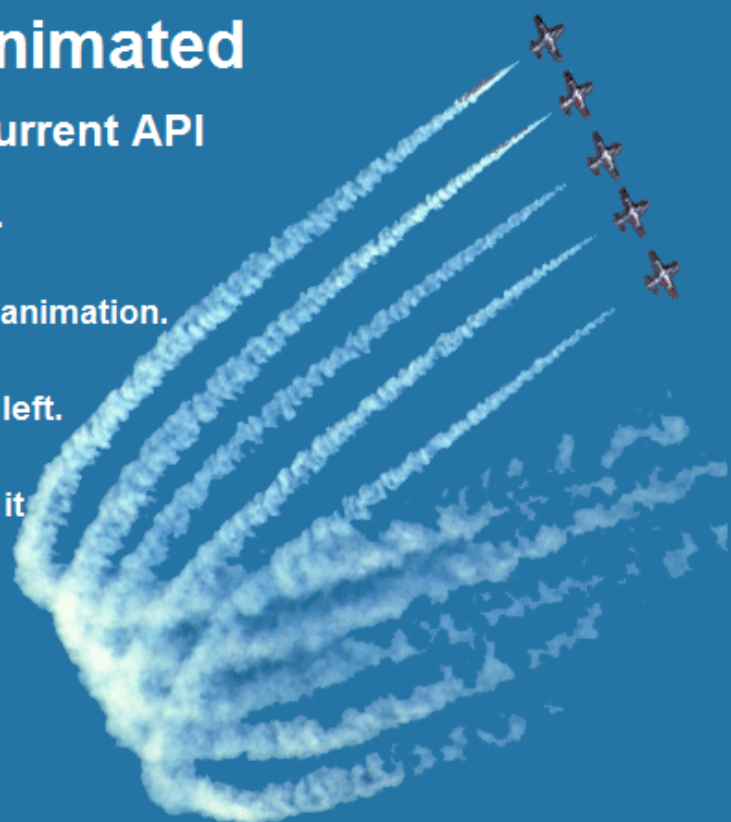
## Java Concurrent Animated

### Visualizing the Java Concurrent API

- Ctrl-R resets the current animation.
- Ctrl-S pauses/resumes the current animation.
- Ctrl-H toggles the menu bar on the left.
- Drag inside snippet panel to resize it
- Use Up and Down arrows to page between the slides.

Written by

- Victor J. Grazi  
 VP Application Development - Credit Suisse  
 vgrazi@gmail.com  
 Download: <http://sourceforge.net/projects/javaconcurrenta/>



# Bibliography

VICTOR J. GRAZI ; (2002) **Java Concurrent Animated**. Application  
<http://sourceforge.net/projects/javaconcurrenta/>

