



OPERATING SYSTEMS

UNIT 2

PROCESS MANAGEMENT

INTRODUCTION

What we will learn ?

- **What does the operating system do to run...**
 - A program?
 - Several programs simultaneously?
- **How does it do it?**
 - Algorithms
 - Data Structures

}

Processes, threads,
Multiprogramming
Timesharing ...

}

Lifecycle
PCB, Process Table ...

CONTENTS

- 1. Process Management basics**
2. Elements of a Process
3. Process lifecycle
4. Threads (or lightweight processes)
5. Process and thread planning
6. System services for process and thread management

THE CONCEPT OF PROCESS

1- PROCESS MANAGEMENT BASICS

- A process is **any running program**.
- A process needs certain **RESOURCES** to perform its task satisfactorily:
 - CPU time.
 - Memory.
 - Files.
 - I / O Devices.
- Resources are allocated to a process:
 - When creating.
 - During execution

THE CONCEPT OF PROCESS

1- PROCESS MANAGEMENT BASICS

- A **process** is alive
- An instant in the process's life is given by
 - Its **code**
 - **Data** that has at that moment
 - The value of the execution **stack**
 - The value of all processor **registers**
 - The **program counter**
 - A pointer to the next instruction to be executed

PROCESS MANAGER FUNCTIONS

1- PROCESS MANAGEMENT BASICS

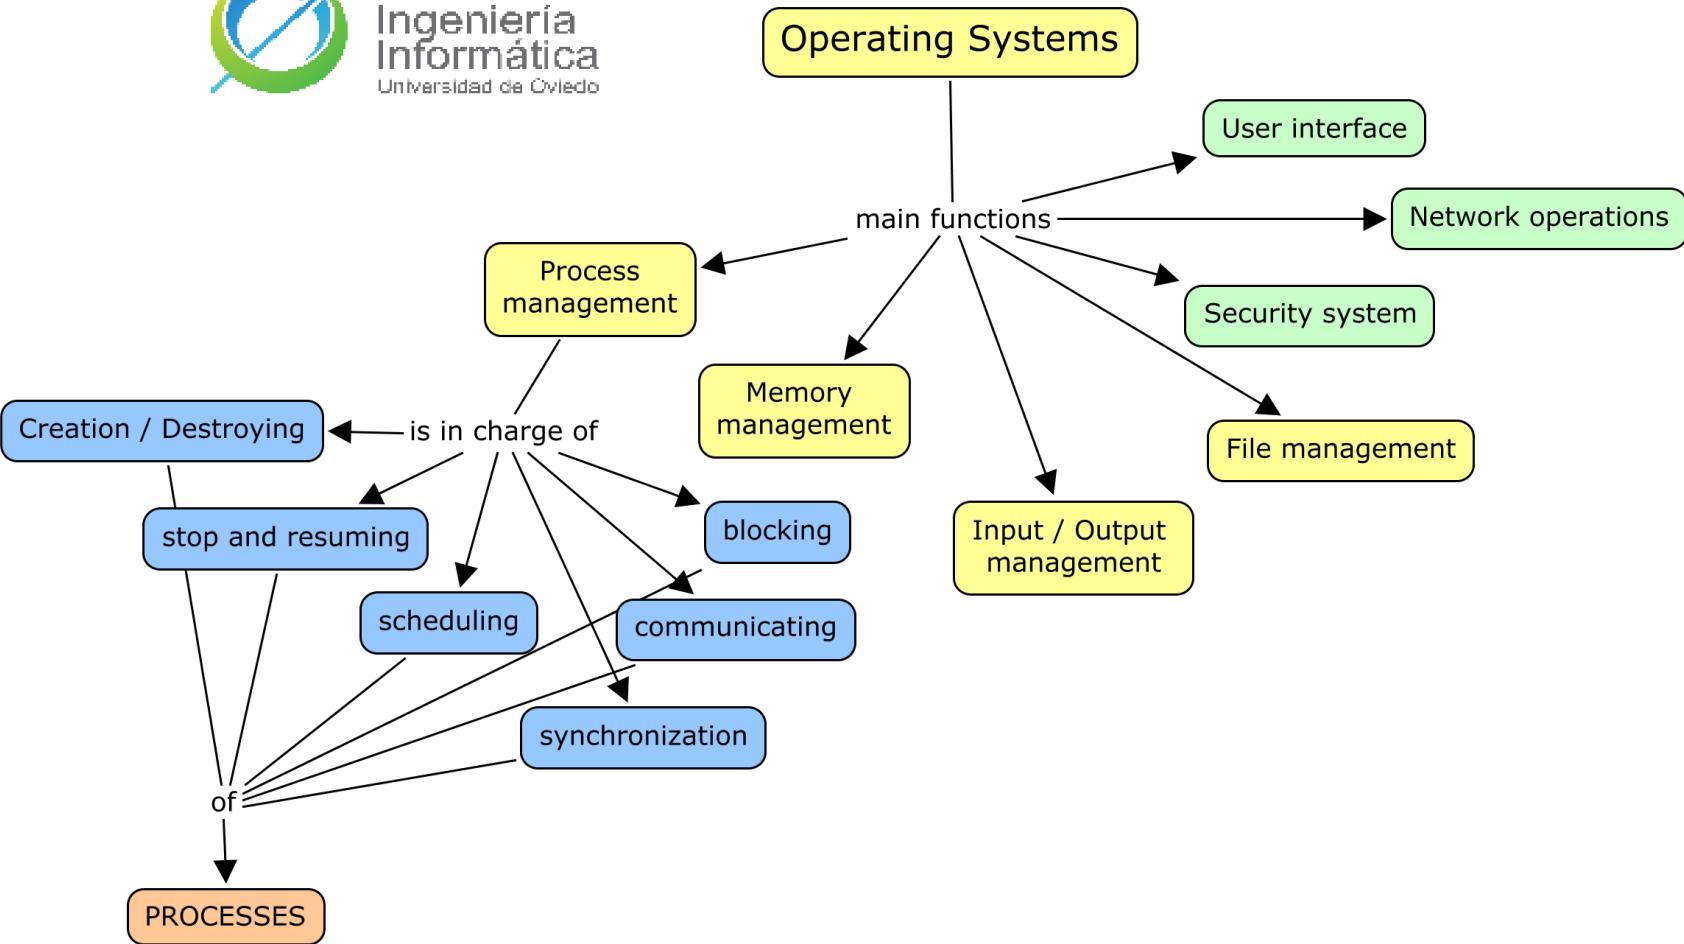
- The obligations of the OS as process manager are:
 - **Creating and deleting** processes.
 - **Process scheduling** (ensuring the execution of multiple processes to maximize processor utilization).
 - Establishing mechanisms for process **synchronization and communication**.
 - Handling **deadlocks**.

PROCESS MANAGER FUNCTIONS

1- PROCESS MANAGEMENT BASICS



Escuela de
Ingeniería
Informática
Universidad de Oviedo



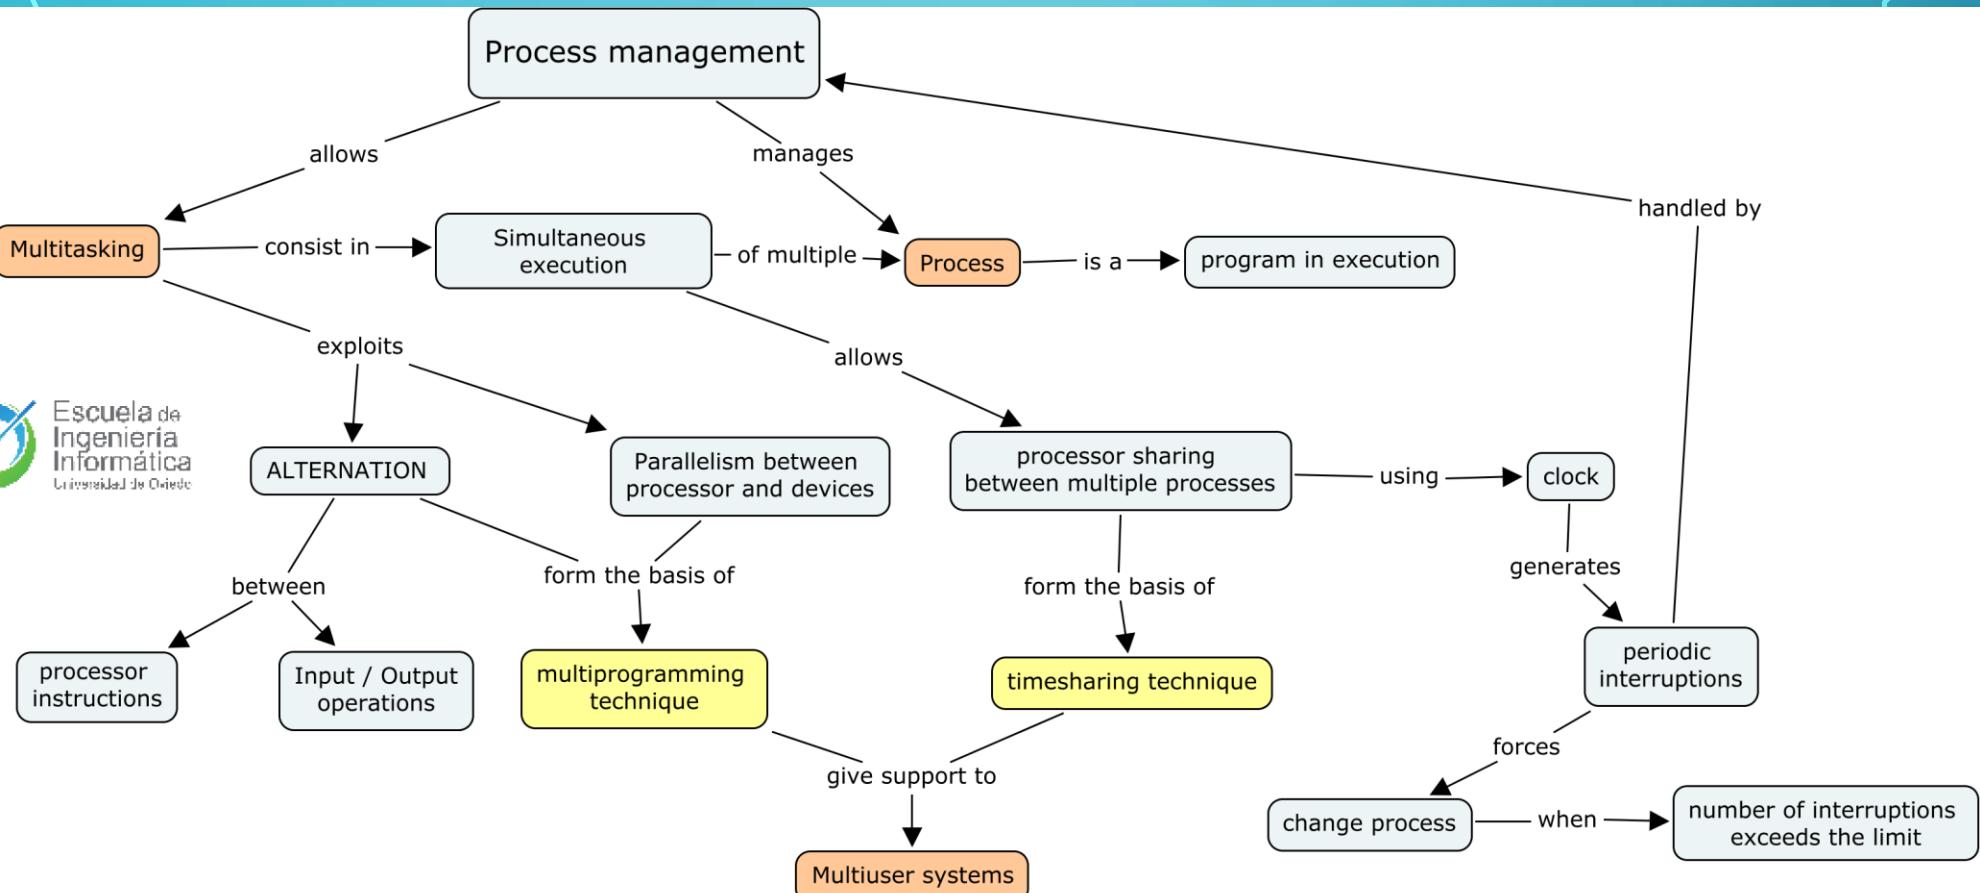
- **How can many processes be executed simultaneously on a machine with a single processor ?**

1. *Parallelism between processor and I / O*
2. *Programs alternation between processor and I / O*
3. *Processor time sharing among processes*



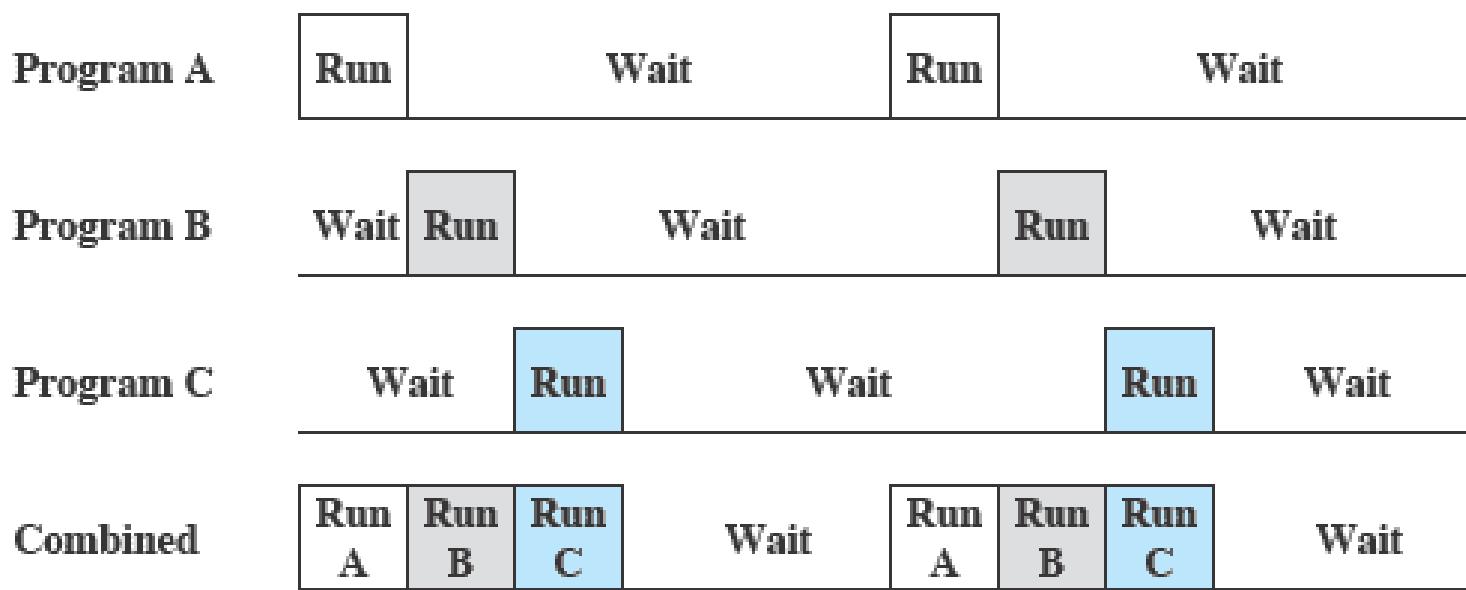
MULTITASKING

1- PROCESS MANAGEMENT BASICS



MULTITASKING

1 - PROCESS MANAGEMENT BASICS



(c) Multiprogramming with three programs

1. Process Management basics
- 2. Elements of a Process**
3. Process lifecycle
4. Threads (or lightweight Processes)
5. Process and thread planning
6. System services for process and thread management

2.- *ELEMENTS OF A PROCESS*

What information does the OS need to handle about a process to run it on Multitasking?



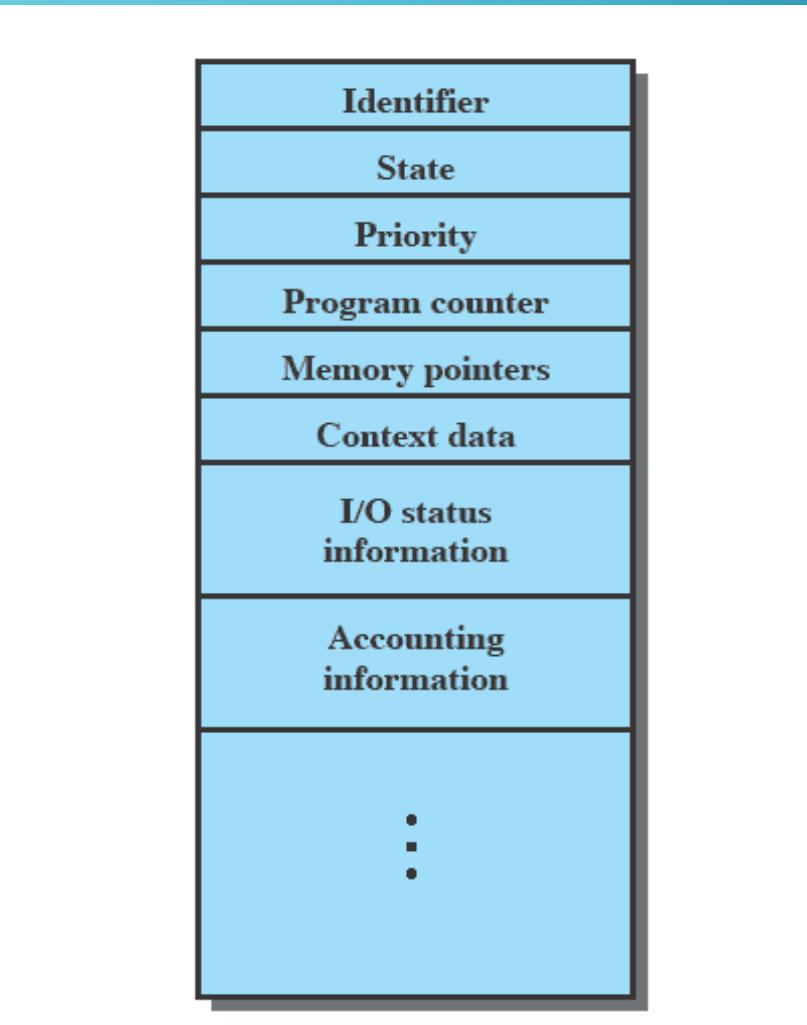
2.- ELEMENTS OF A PROCESS

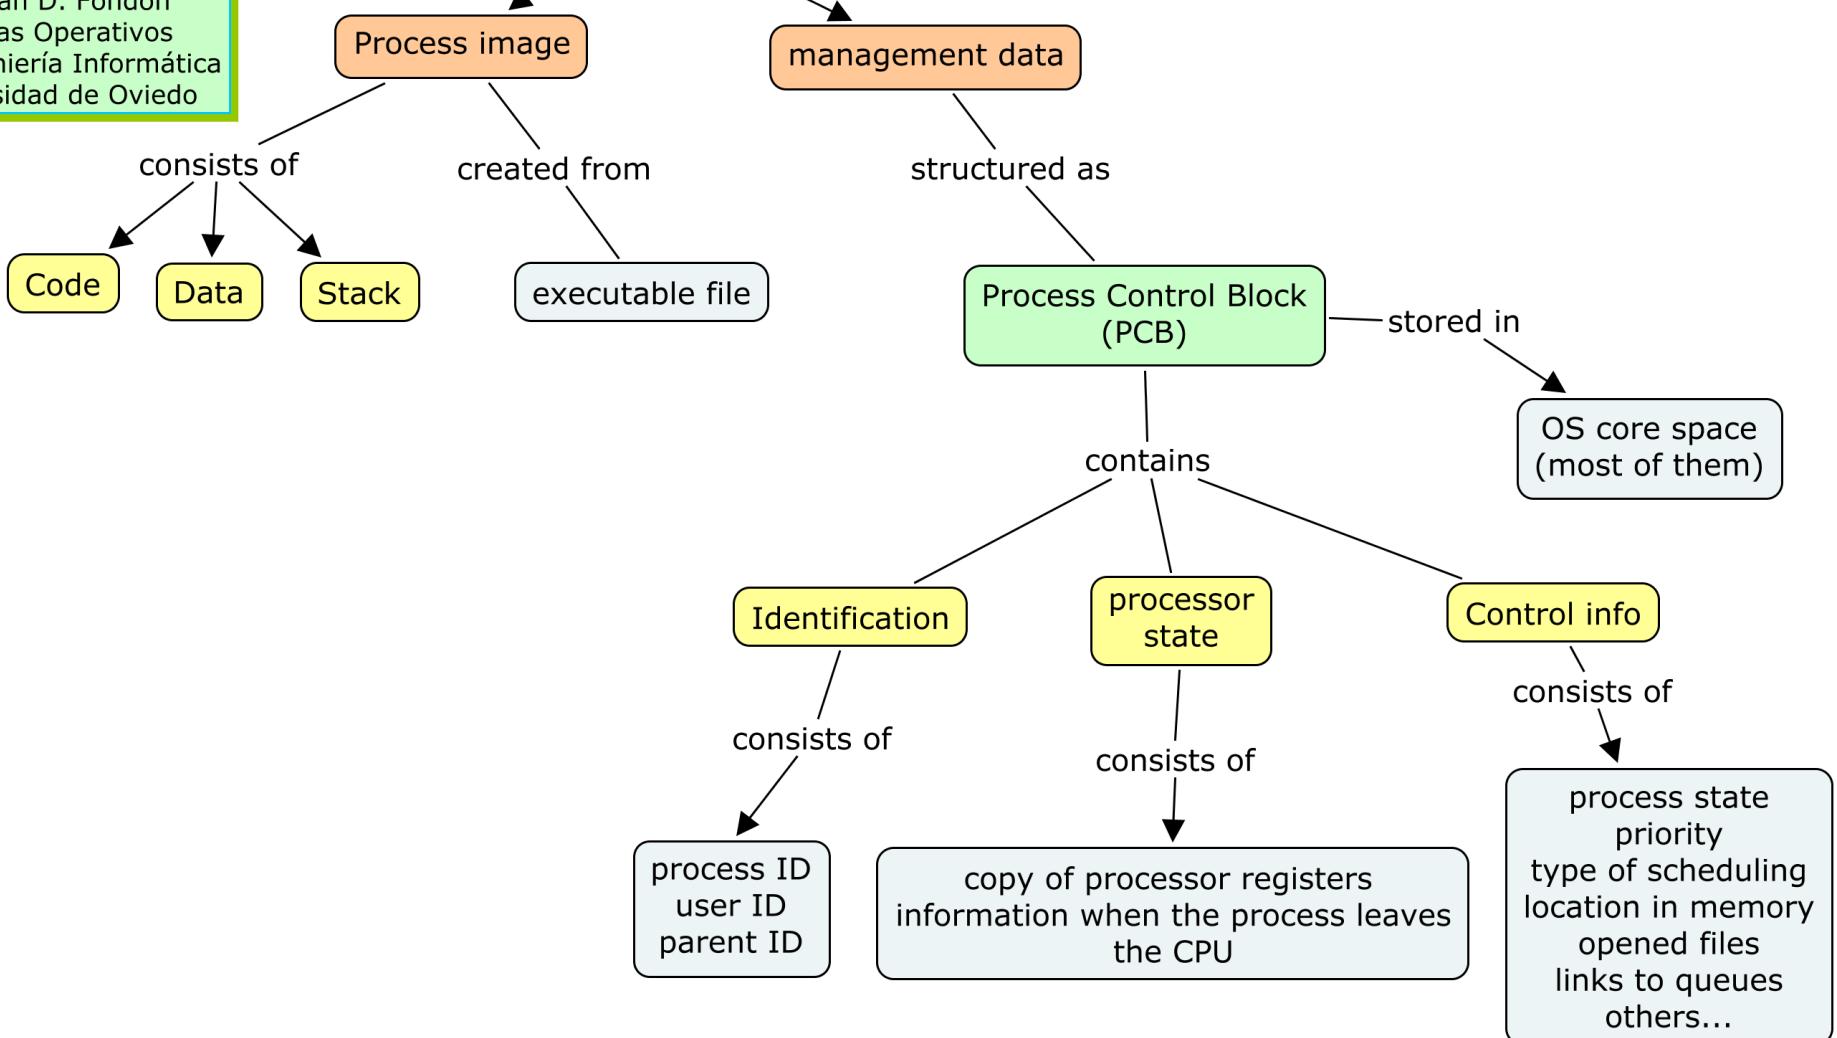
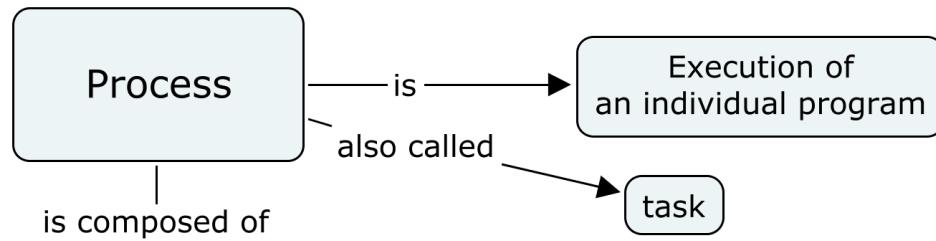
- A process is:
 - **Text section** (program code)
 - **Data section** containing *global variables* and *dynamic memory*
 - **Stack**, which contains *temporary data* (subroutine parameters, return addresses and local variables)
 - Current activity, represented in a **PCB** data structure:
 - Value of the **program counter**
 - Contents of **processor registers**
 - Priority, status, open files, memory space location, ...
 - The OS usually uses a **Table of PCBs** to locate all processes in the system.

Process Control Block

2.- Elements of a Process

- Data structure that contains the process elements
- It is possible to interrupt a running process and later resume execution as if the interruption had not occurred
- Created and managed by the operating system
- Key tool that allows support for multiple processes



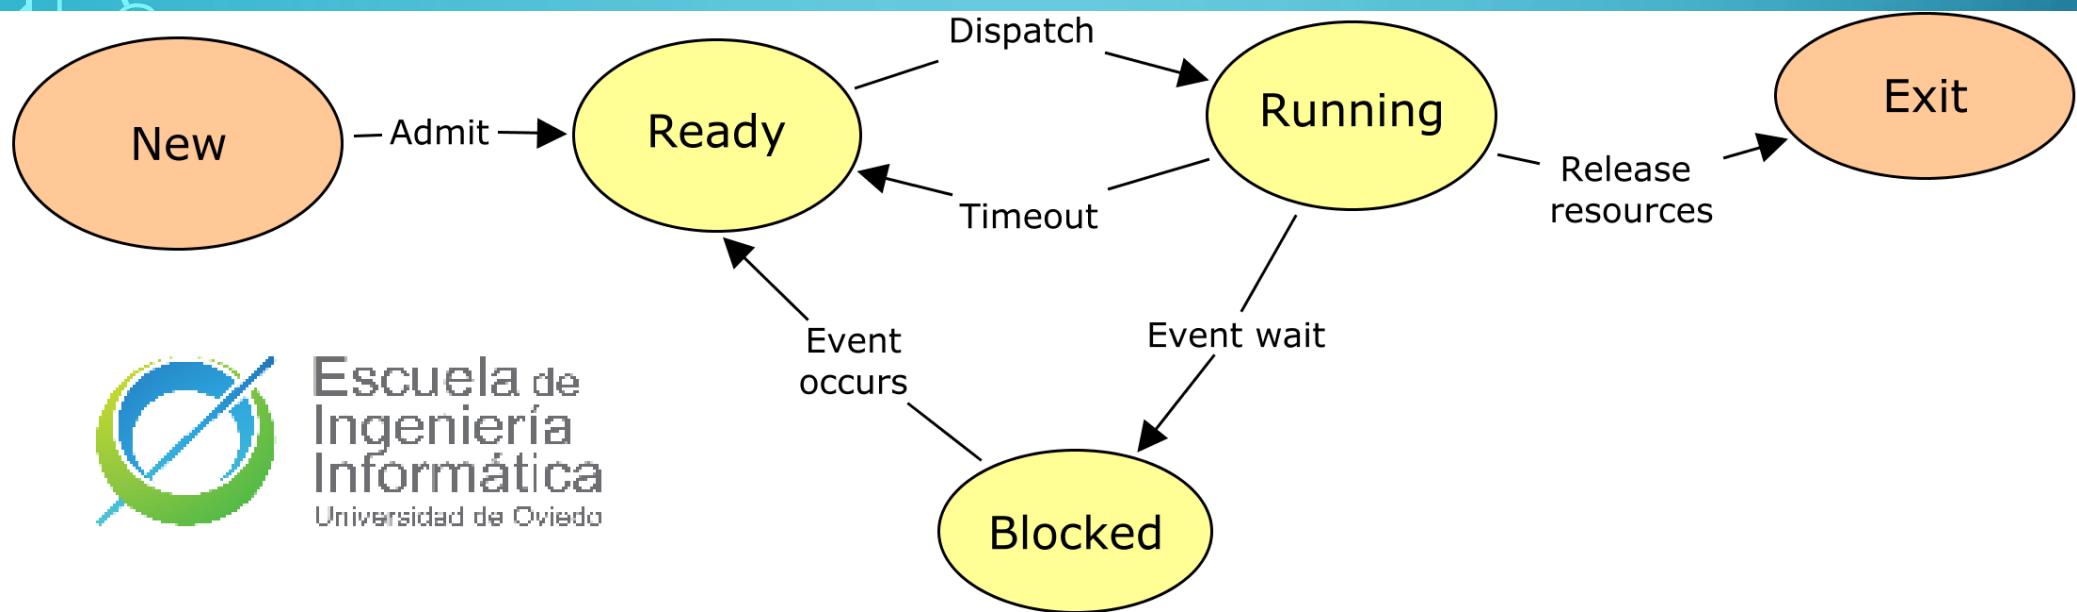


CONTENTS

1. Process Management basics
2. Elements of a Process
- 3. Process lifecycle**
4. Threads
5. Process and thread planning
6. System services for process and thread management

FIVE STATE MODEL

3.- PROCESS LIFECYCLE

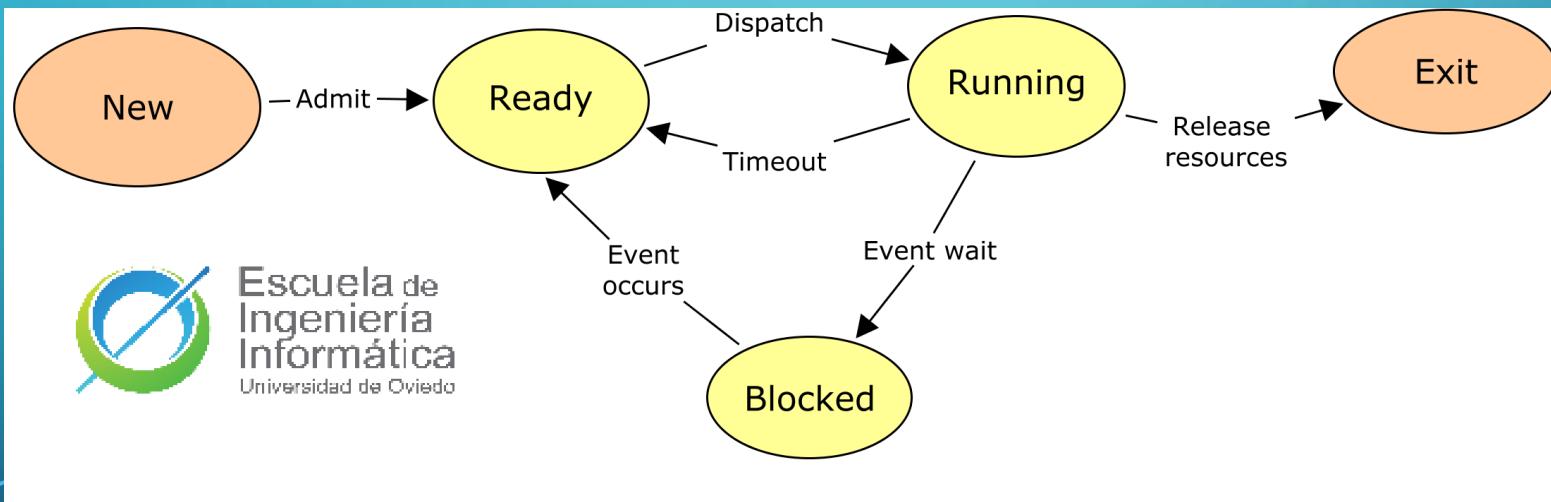


Escuela de
Ingeniería
Informática
Universidad de Oviedo

FIVE STATE MODEL

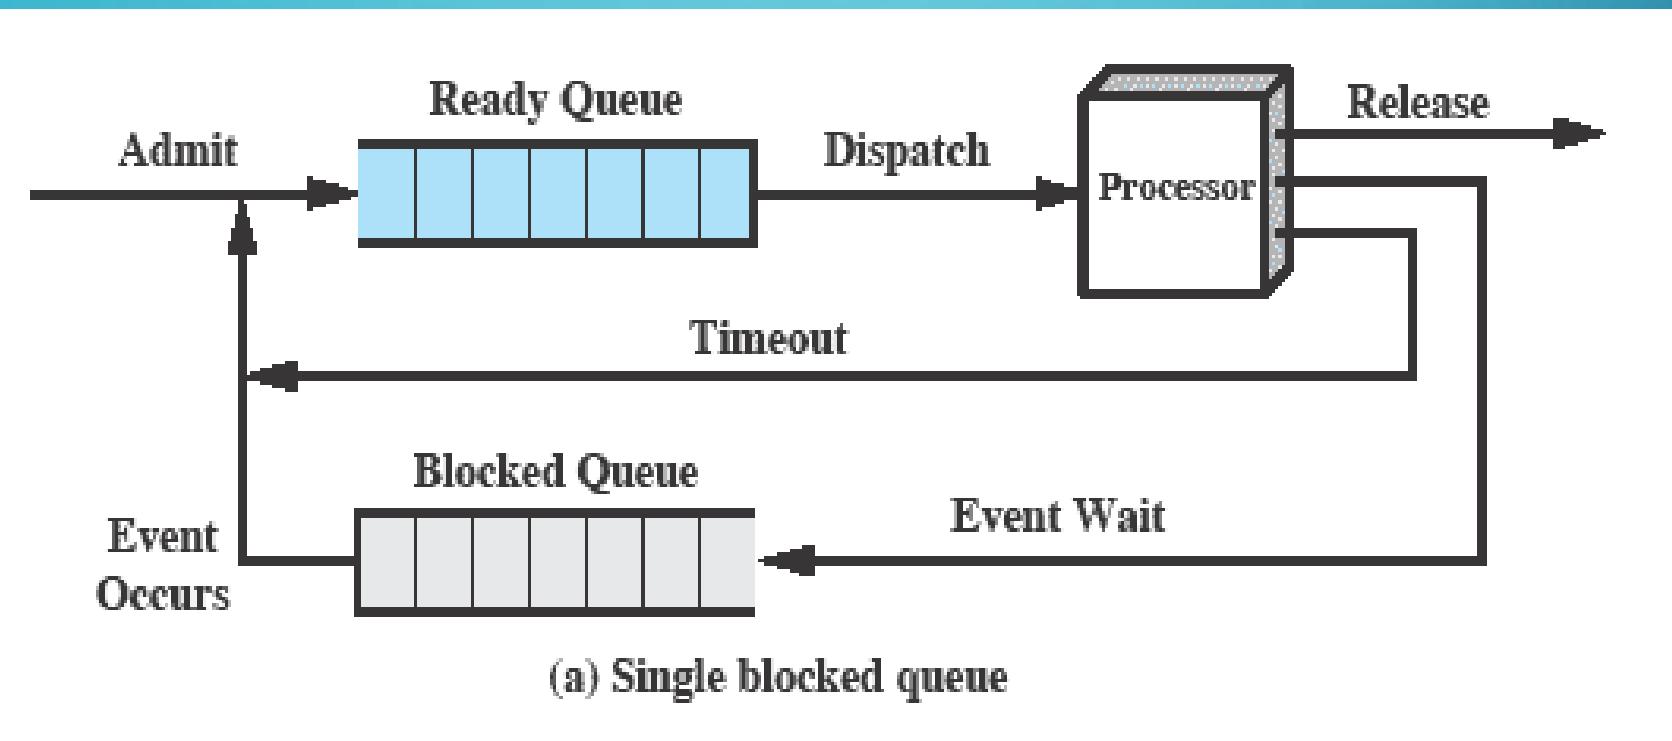
3- PROCESS LIFECYCLE

1. What situations does imply a state change ?



FIVE STATE MODEL

3- PROCESS LIFECYCLE



FIVE STATE MODEL

3- PROCESS LIFECYCLE

- As a process is executed, its **state changes**. Each process can be in one of this states:
 - **New:** The process is being created.
 - **Running:** The process is in the CPU executing instructions.
 - **Blocked:** The process is waiting for an event to occur (e.g. I/O completion or signal reception).
 - **Ready:** waiting to be assigned to a processor.
 - **Finished:** execution completed, thus no further instructions to execute and the OS will withdraw the resources it consumes.

FIVE STATE MODEL

3- PROCESS LIFECYCLE

- State transitions in this model are:
 - **None to New:** creates a new process to run a program
 - **New to Ready:** the system is ready to accept a new process as it has enough resources for it.
 - **Ready to Running:** the system chooses one of the processes in the ready state to carry out execution.
 - **Running to Finished:** the running process is terminated by the OS if it indicates that it has ended, abandoned or canceled.
 - **Running to Ready:** the process has exhausted its running time, or yields voluntarily or it is stopped to execute other process (possibly with higher priority).

FIVE STATE MODEL

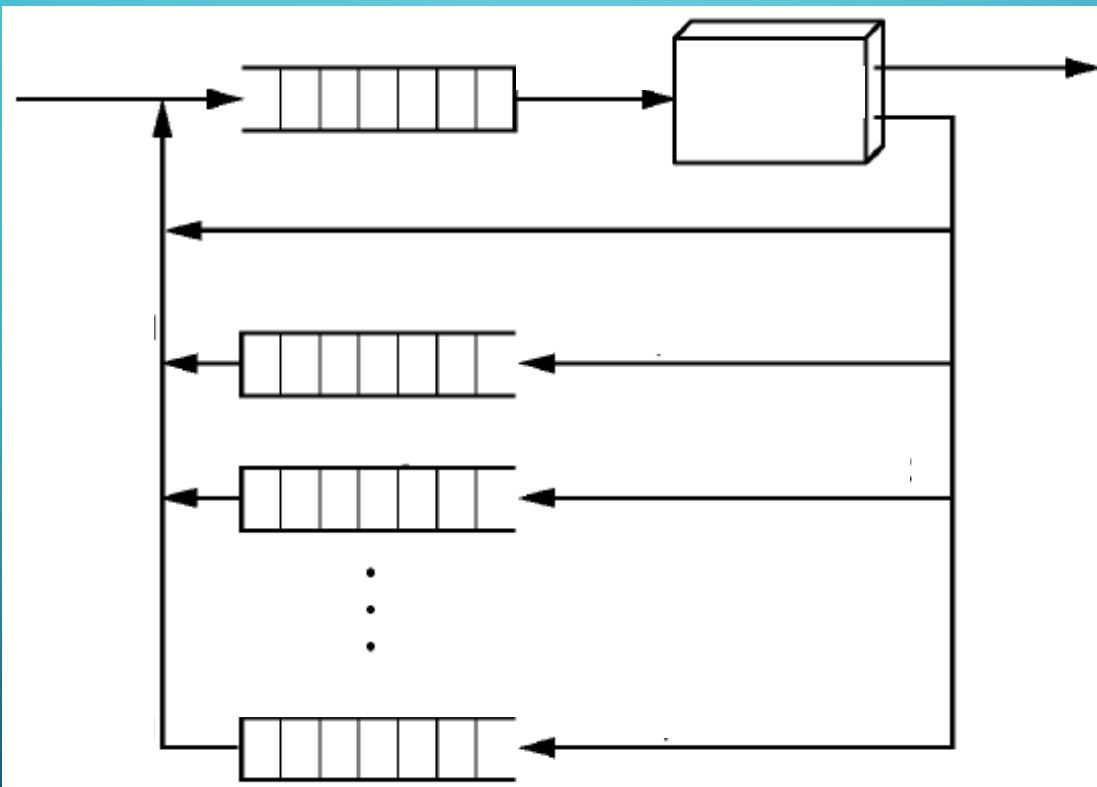
3- PROCESS LIFECYCLE

- State transitions 2nd part:
 - **Running to Blocked:** the process asks for a (slow) operation and waits for it
 - **Blocked to Ready:** the (slow) operation the process is waiting for eventually finishes.
 - **Ready to Finished** (not shown in the figure): a parent may terminate a child process at any time, or if the parent is completed all children can be completed.
 - **Blocked to Finished:** the same criteria as above.

IMPLEMENTATION

3- PROCESS LIFECYCLE

The states have queues associated for the processes that are in that state



CONTEXT SWITCHING

3- PROCESS LIFECYCLE

Context Switching

Operations performed by the OS to make a process release the processor, and make another one obtain the processor.

Implies:

- Saving processor state in the PCB for the process that leaves the CPU.
- Restoring the processor state from the process's PCB that is assigned to the CPU.

OS INTERRUPTION HANDLING

3- PROCESS LIFECYCLE

Process management issues

- What happens when an interruption occurs? How do you jump to the OS code?
- What does a process change imply?
- What kinds of interruption necessarily cause a process change and which don't?
- What does the creation of a process imply?



OS interruption handling

3- Process lifecycle

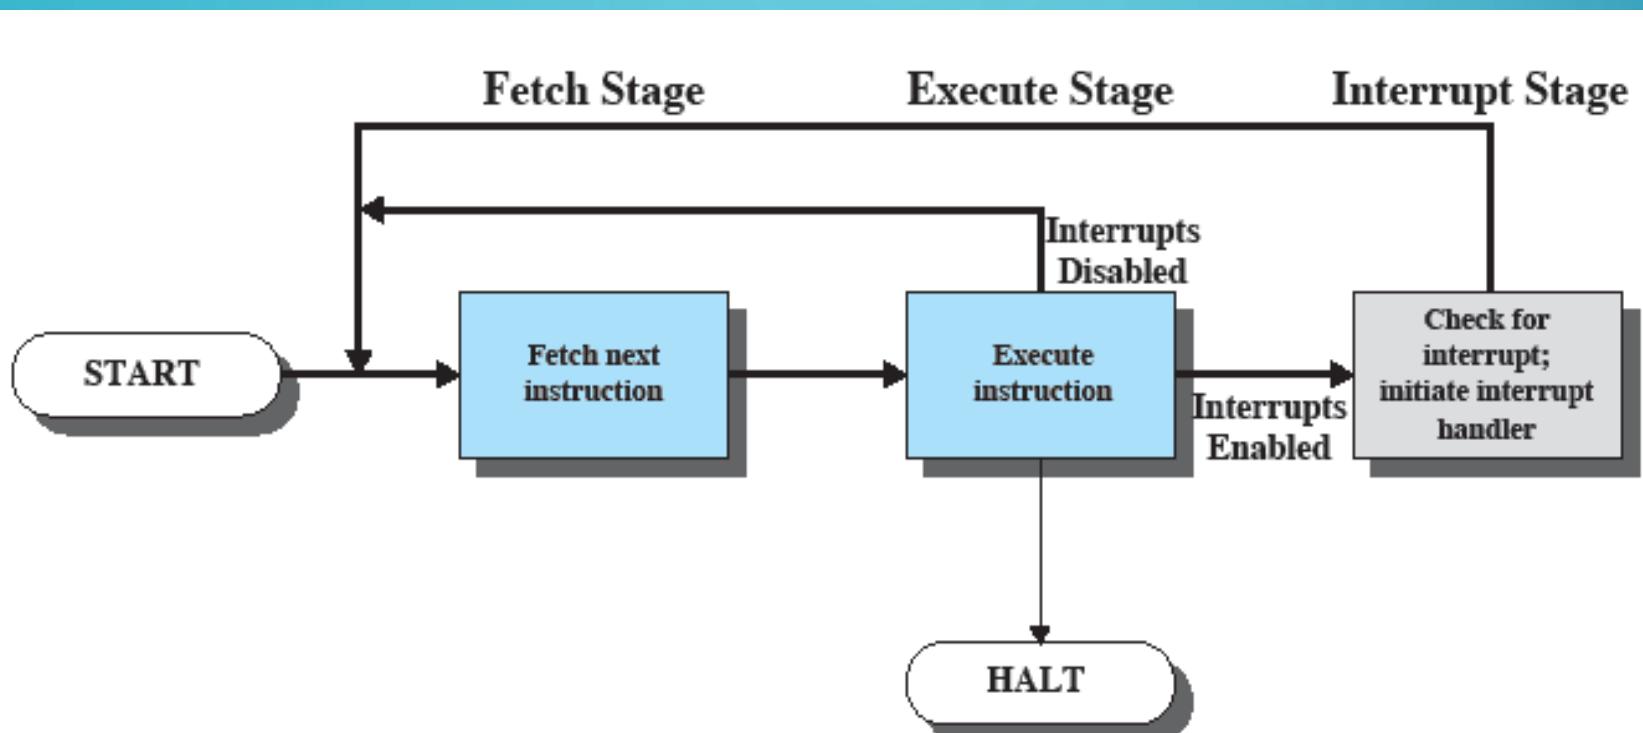


Figure 1.7 Instruction Cycle with Interrupts

INTRODUCTION

What will we learn ?

- **Process an interruption in detail**

- A **very simple** interruption
- It doesn't involve any process switching
- The OS just does some **very small task** and goes back to the program
- The OS is in the same address space as the program (!?)

INTERRUPTION HANDLING

1 - IRQ number 5 arrives

- **Tasks for Figure 1:**

- Draw the lines (pointers) from PC, BSP, SP to the memory
- Figure out at which **memory address** the handling routine for IRQ 5 (i.e., the **code to be executed**) starts

- Note: IRQ means “Interruption ReQuest”
- Note: First bit in PSW is mode (0=user, 1=kernel or privileged)

INTERRUPTION HANDLING

2 - Hardware, in only one operation, a) stores PC and PSW at the bottom of the “control stack”, b) activates kernel mode, c) points PC to first instruction in the handling routine code

- **Tasks for Figure 2:**

- Update numbers in the picture, both processor and memory
- Then draw pointers as needed

INTERRUPTION HANDLING

3 – The rest of the processor registers are stored in control stack, and stack pointers (SP and BSP) point to that stack.

- **Tasks for Figure 3:**

- Update numbers in the picture, both processor and memory
- Then draw pointers as needed

INTERRUPTION HANDLING

- 4—Instructions of interruption handler are executed, until just before the return (IRET is located at memory address 13999).

- **Tasks for Figure 4:**

- Update numbers and draw pointers as usual

INTERRUPTION HANDLING

- 5 – Last instructions of the Interrupt Handler restores general registers and stack registers to the processor

- **Tasks for Figure 5:**

- Update numbers and draw pointers as usual

INTERRUPTION HANDLING

6 – As the last step of “IRET”, hardware restores PC and PSW.

- **Tasks for Figure 6:**

- Update numbers and draw pointers as usual

- After that... is the processor in user mode or in kernel mode?

OS INTERRUPTION HANDLING

3- PROCESS LIFECYCLE

Process management issues

What happens when an interruption occurs? How do you jump to the OS code?

What does imply a process change?

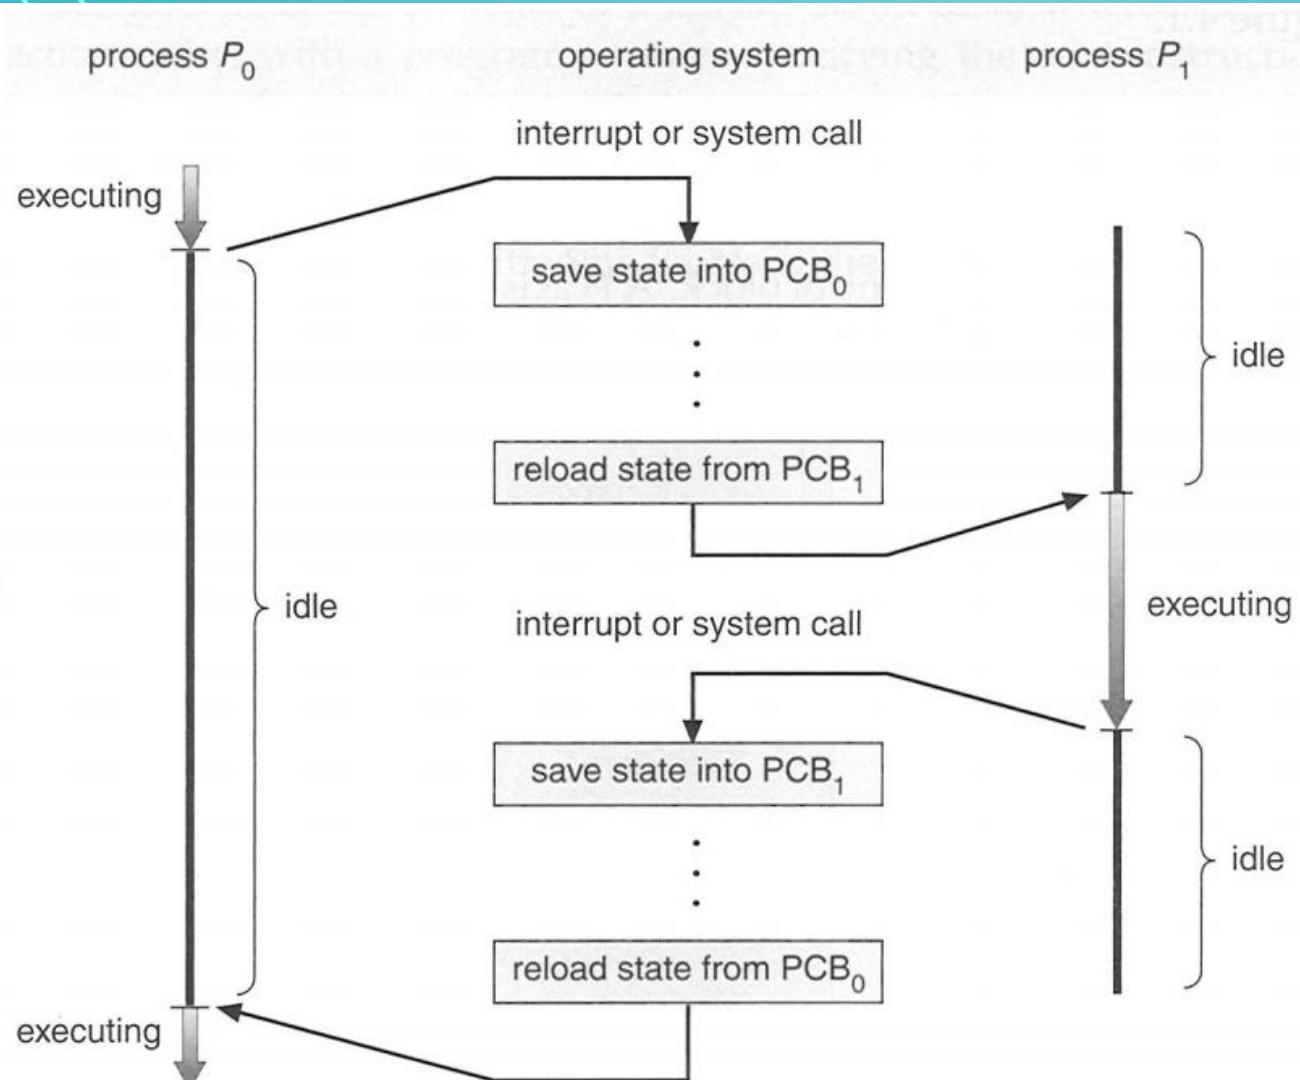
What kinds of interruption necessarily cause a process change and which don't?

What does imply the creation of a process?

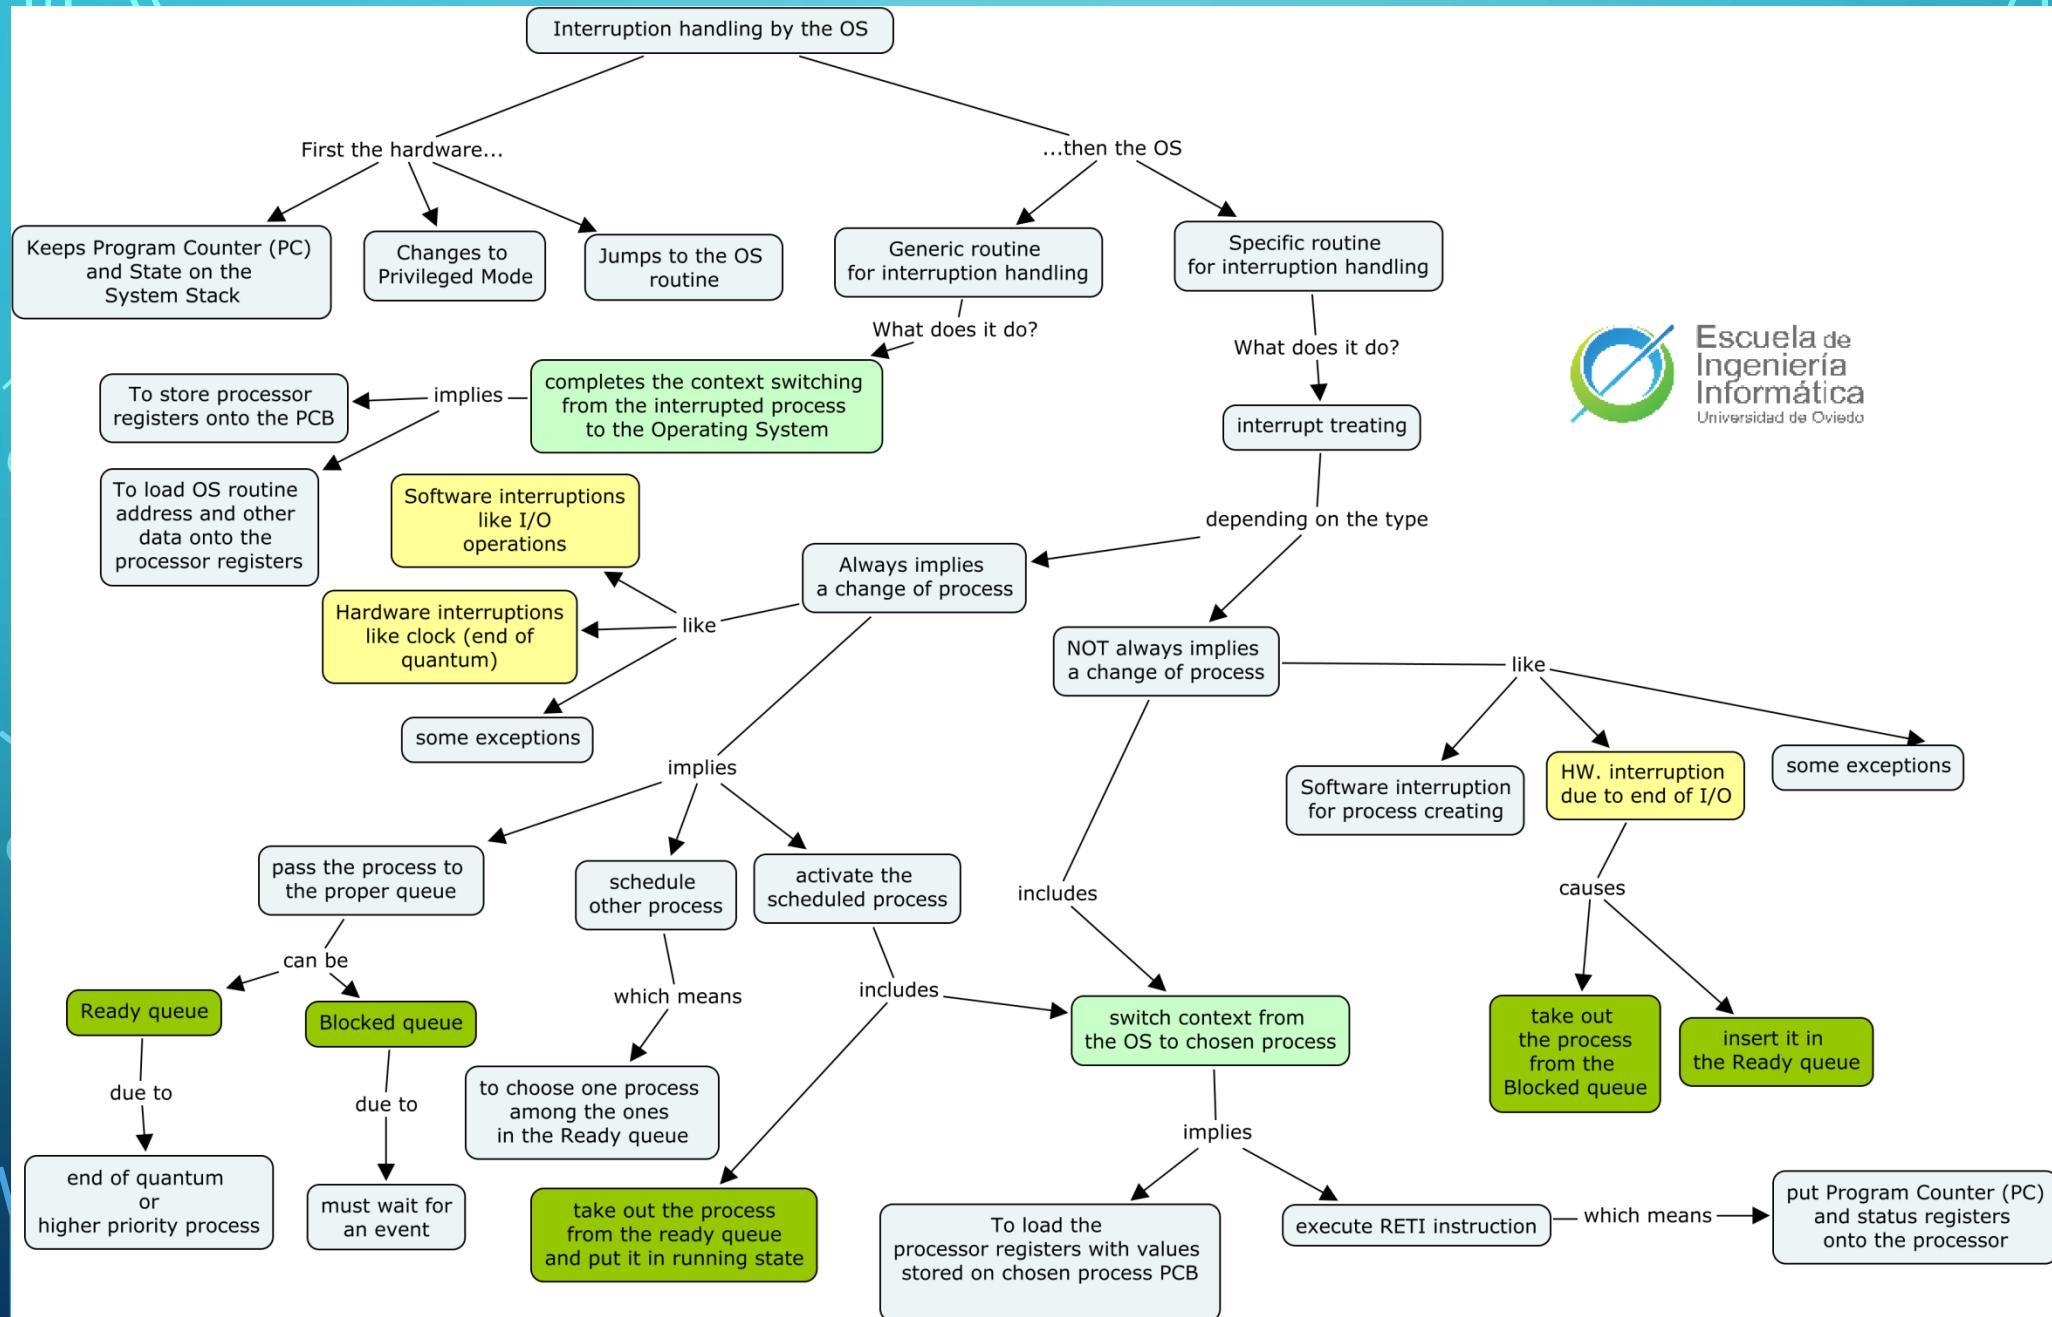


Context Switching

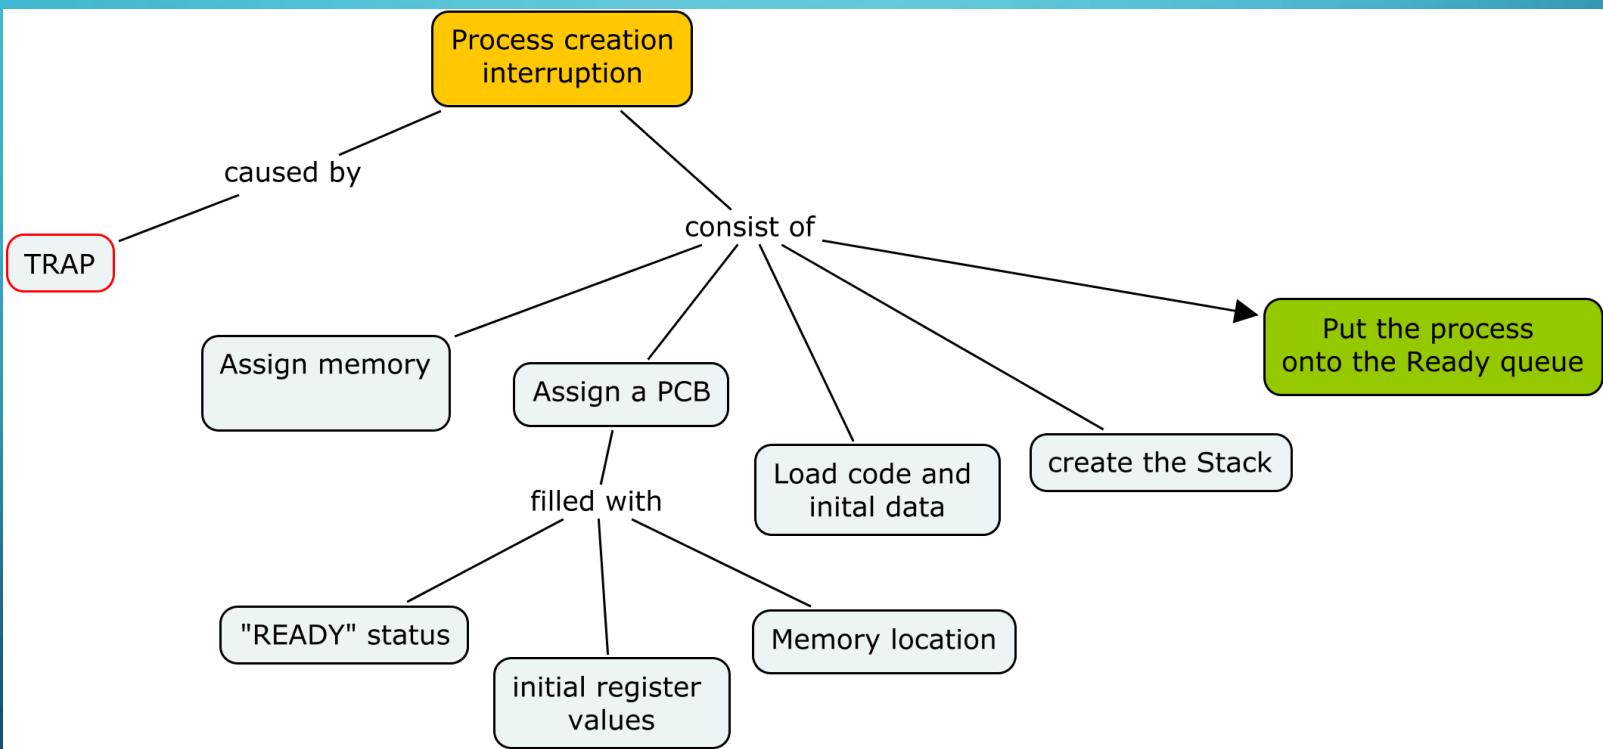
3- Process lifecycle



What does imply a process change?

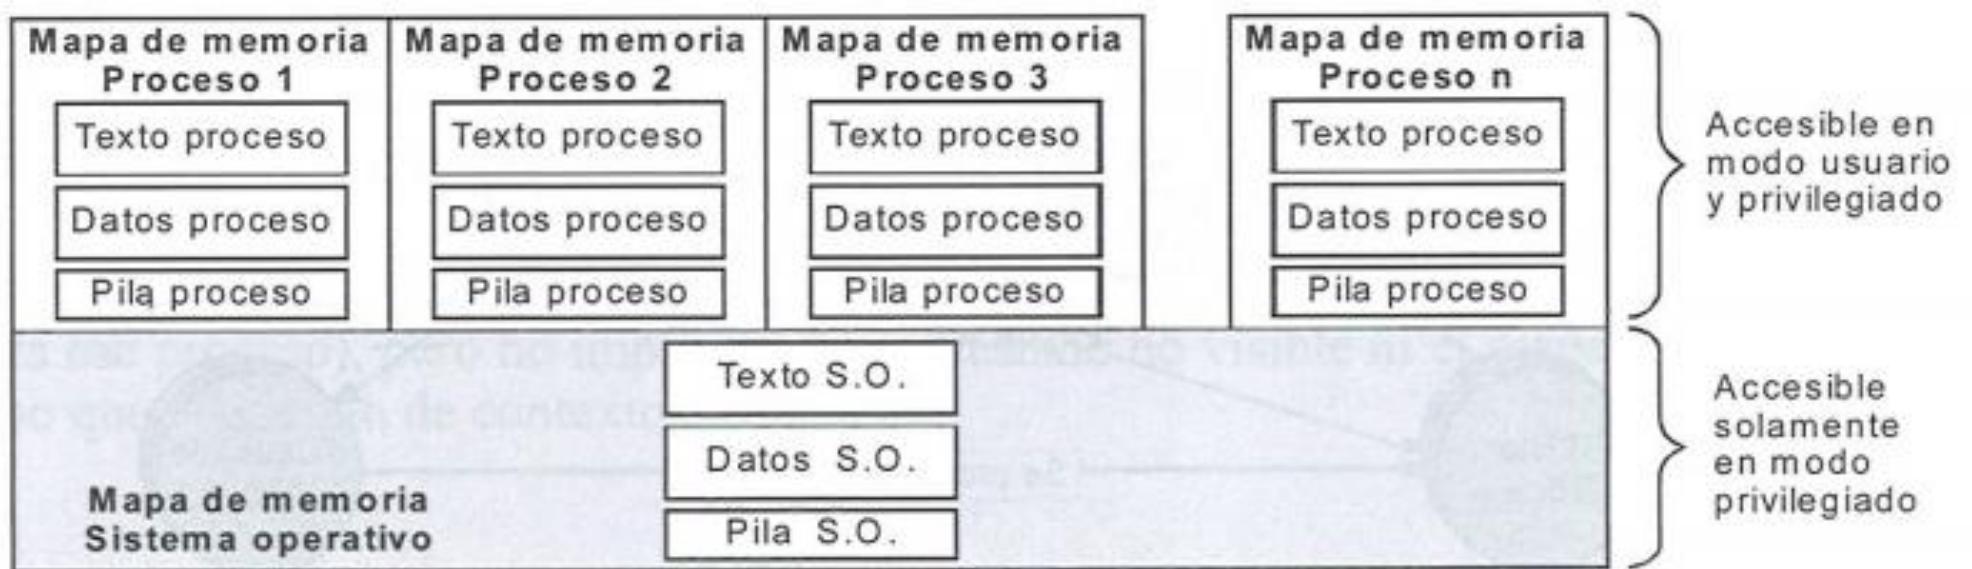


What does imply the creation of a process?



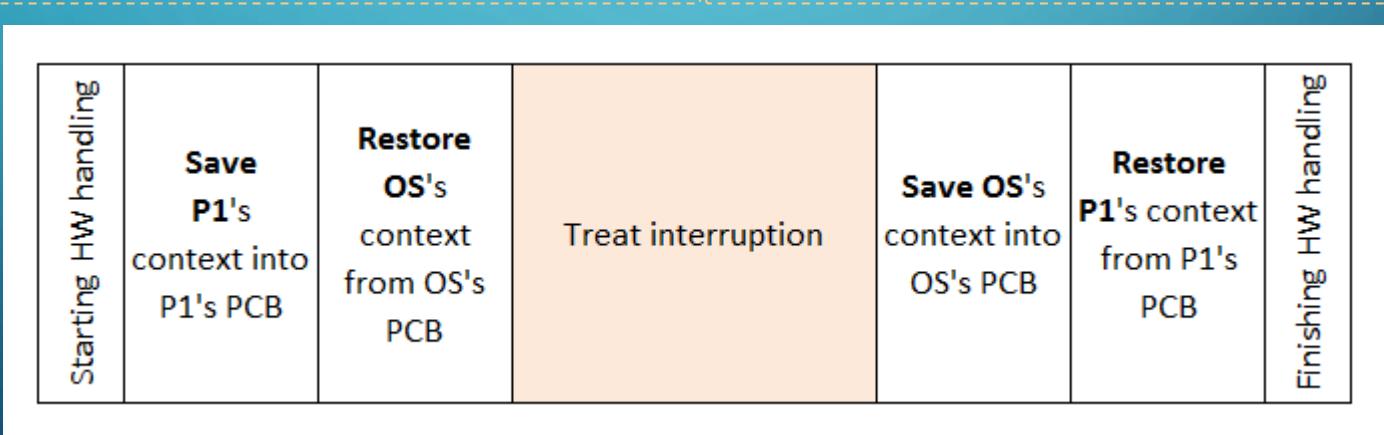
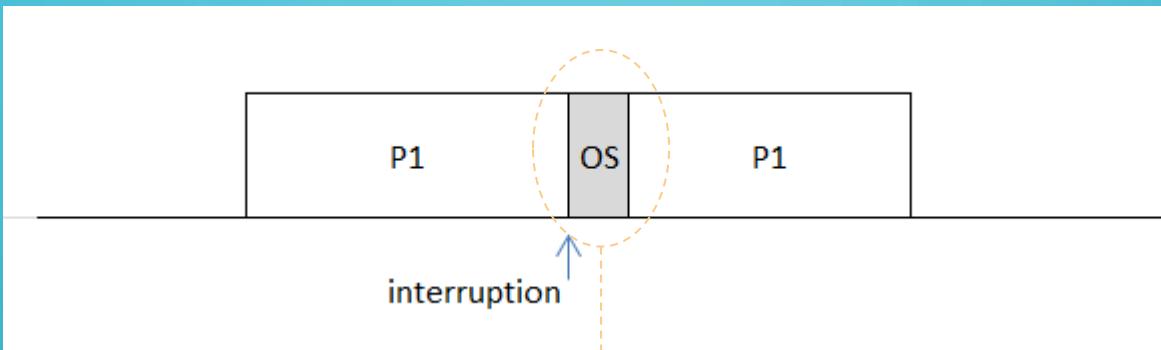
- Independent Kernel

- Performs two context switching on every interruption
- Implies to change the processor mode and the memory map
 - Memory map can be expensive to change
- Can be improved?



Independent kernel execution

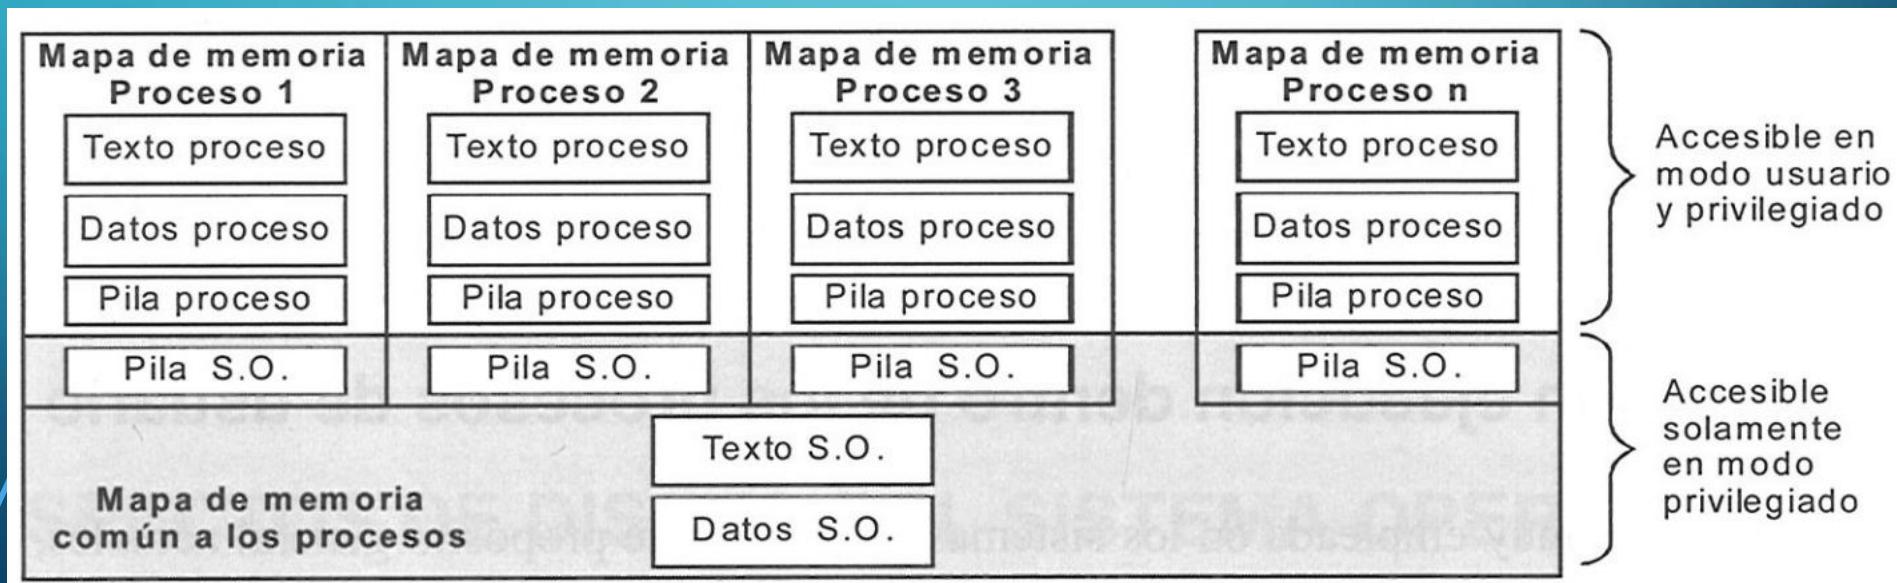
3- Process lifecycle



In-process execution

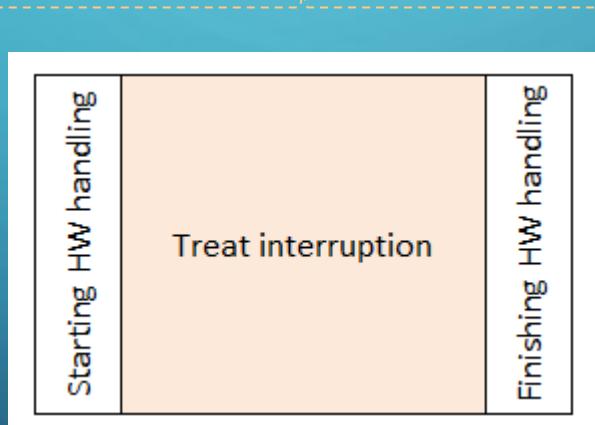
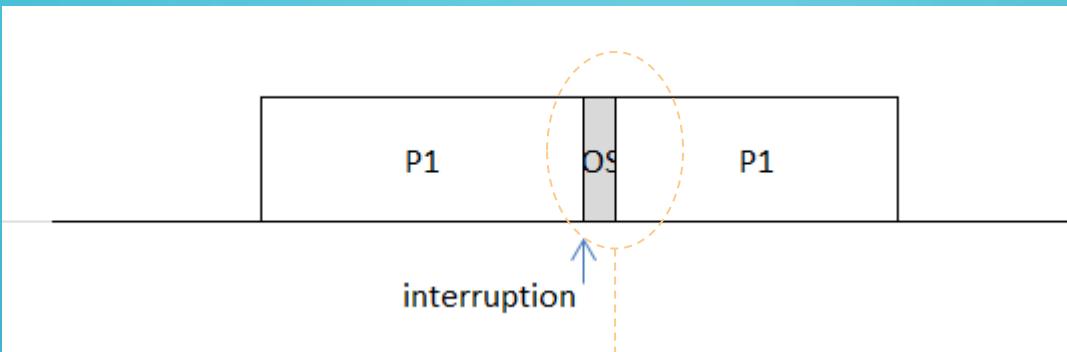
3- Process lifecycle

- In-process execution can avoid context switching
 - Just changes the processor mode



In-process execution

3- Process lifecycle



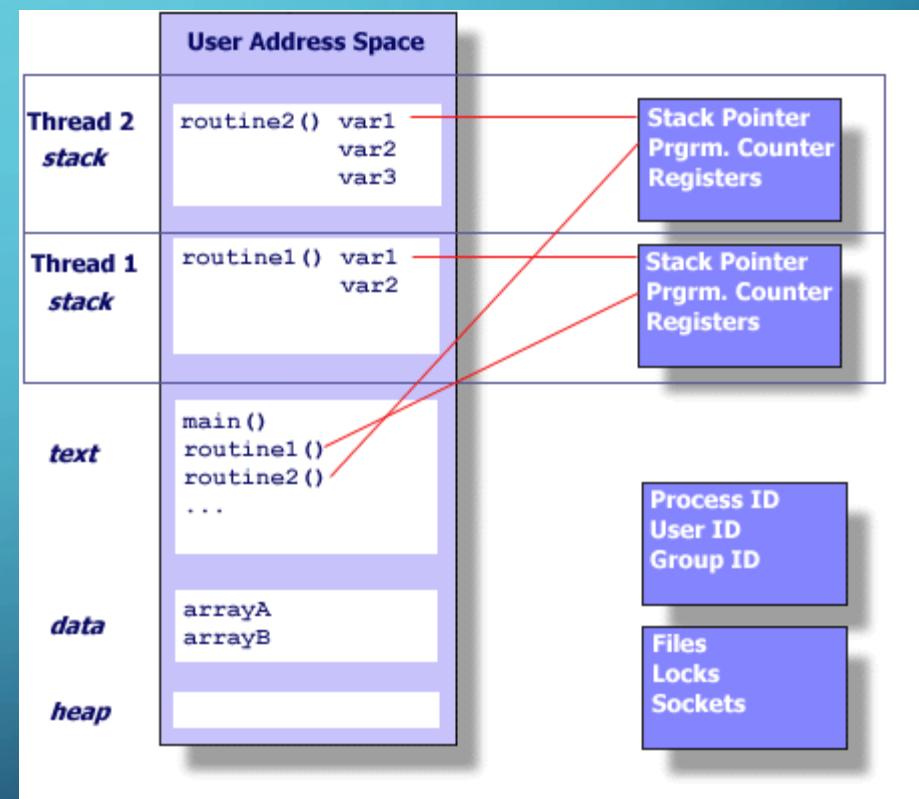
CONTENTS

1. Process Management basics
2. Elements of a Process
3. Process lifecycle
- 4. Threads**
5. Process and thread planning
6. System services for process and thread management

4- THREADS

A flow of execution within a process

- Multiple threads in the same task
- Each thread has:
 - Program Counter,
 - Values of processor registers,
 - Stack
- Threads share the
 - Code (*text*), data and heap
 - Open files, signals, etc..



4- THREADS

Applications

- Parallel Programming (task separation, modularity)

- Information Servers

Advantages

- Ease communication between threads.
- Increase the speed of execution.
- Switching between threads of the same process has little cost.
- The cost of creation and destruction of threads is much lower than with processes
- It therefore improves performance.

Disadvantages

- Sharing memory space (concurrent access to shared resources)
- Greater programming difficulty

Applications... for example?

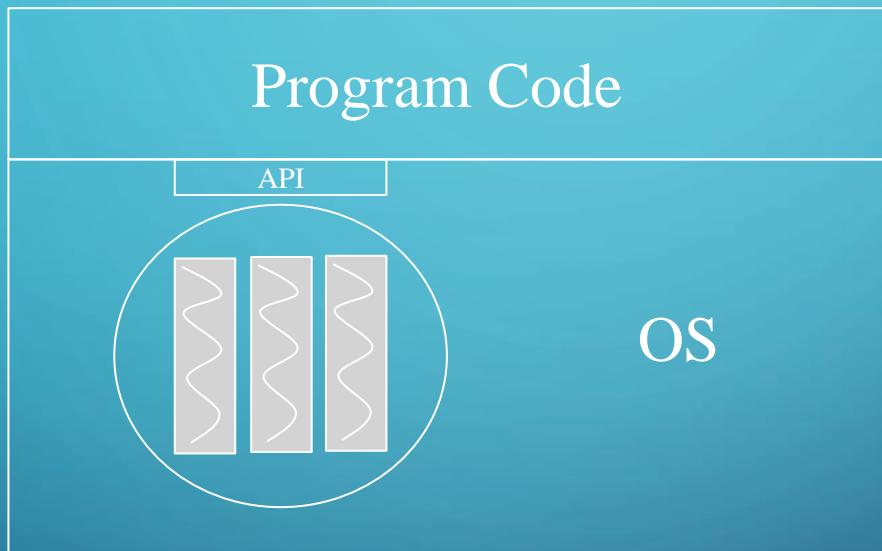
Advantages... Why?



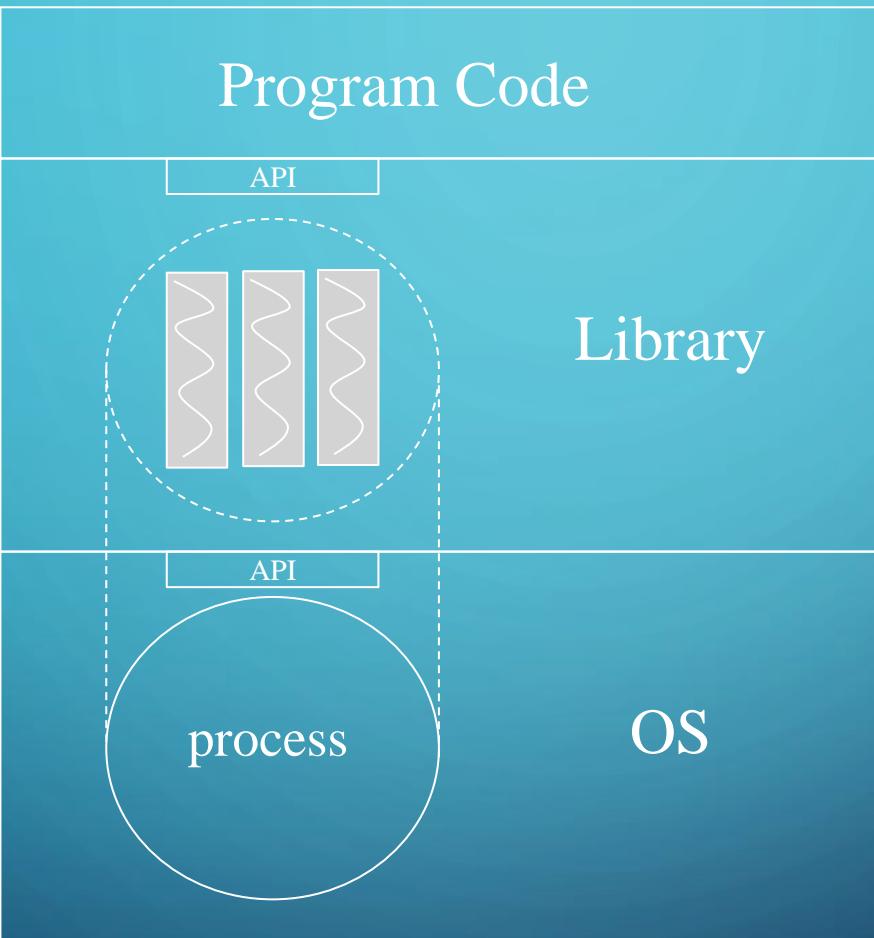
Implementation

- At kernel level (KLT)
- At user level (ULT)
- Combination of both

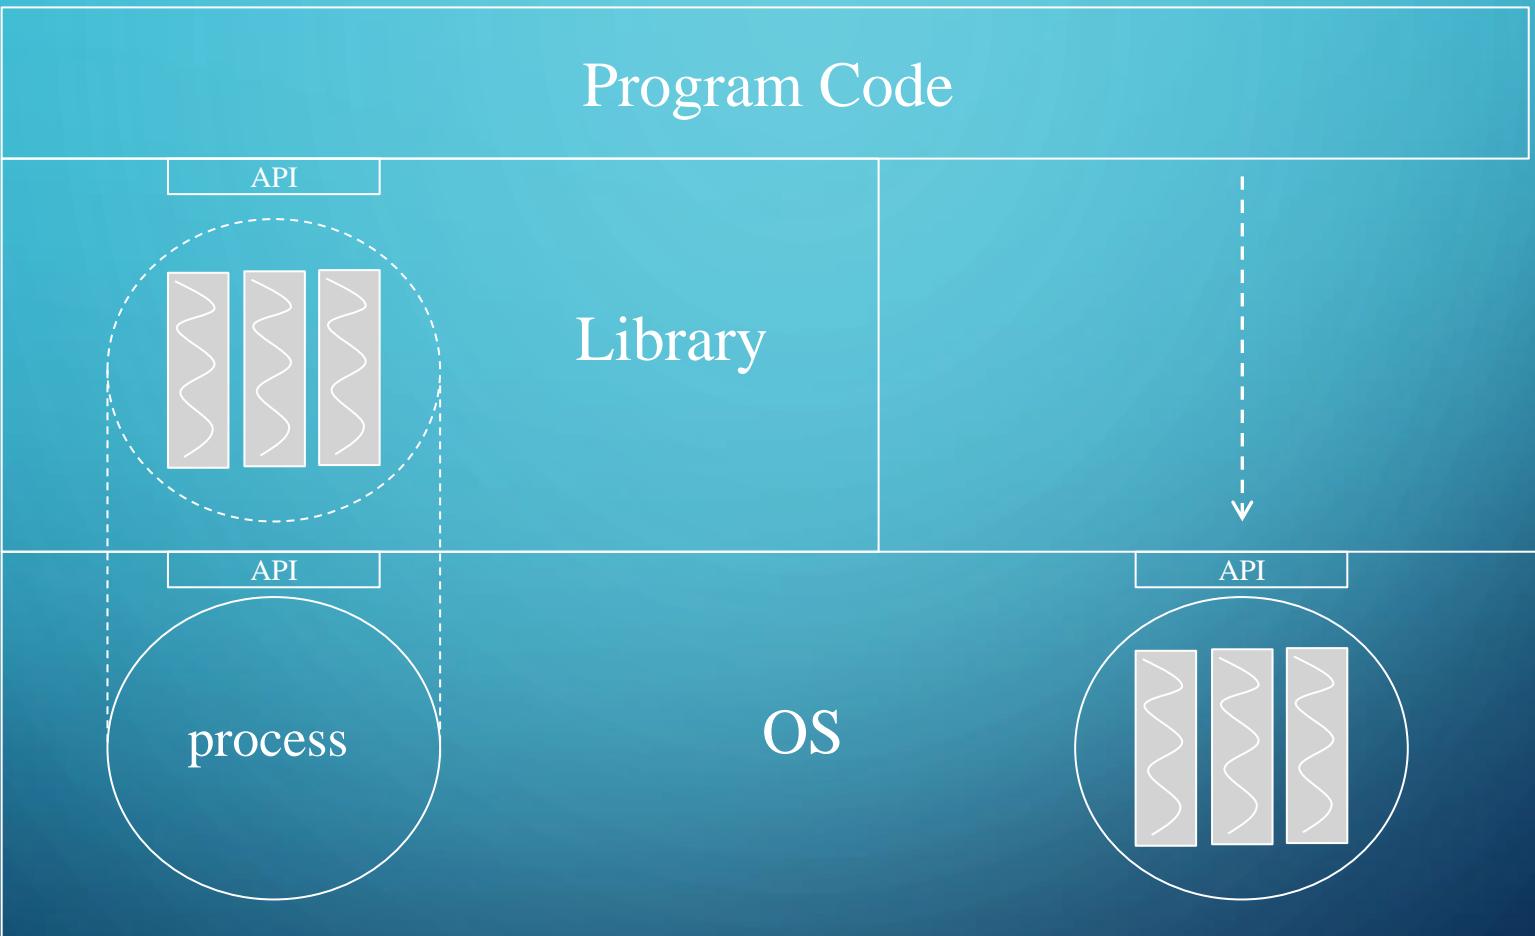
Kernel Level Implementation



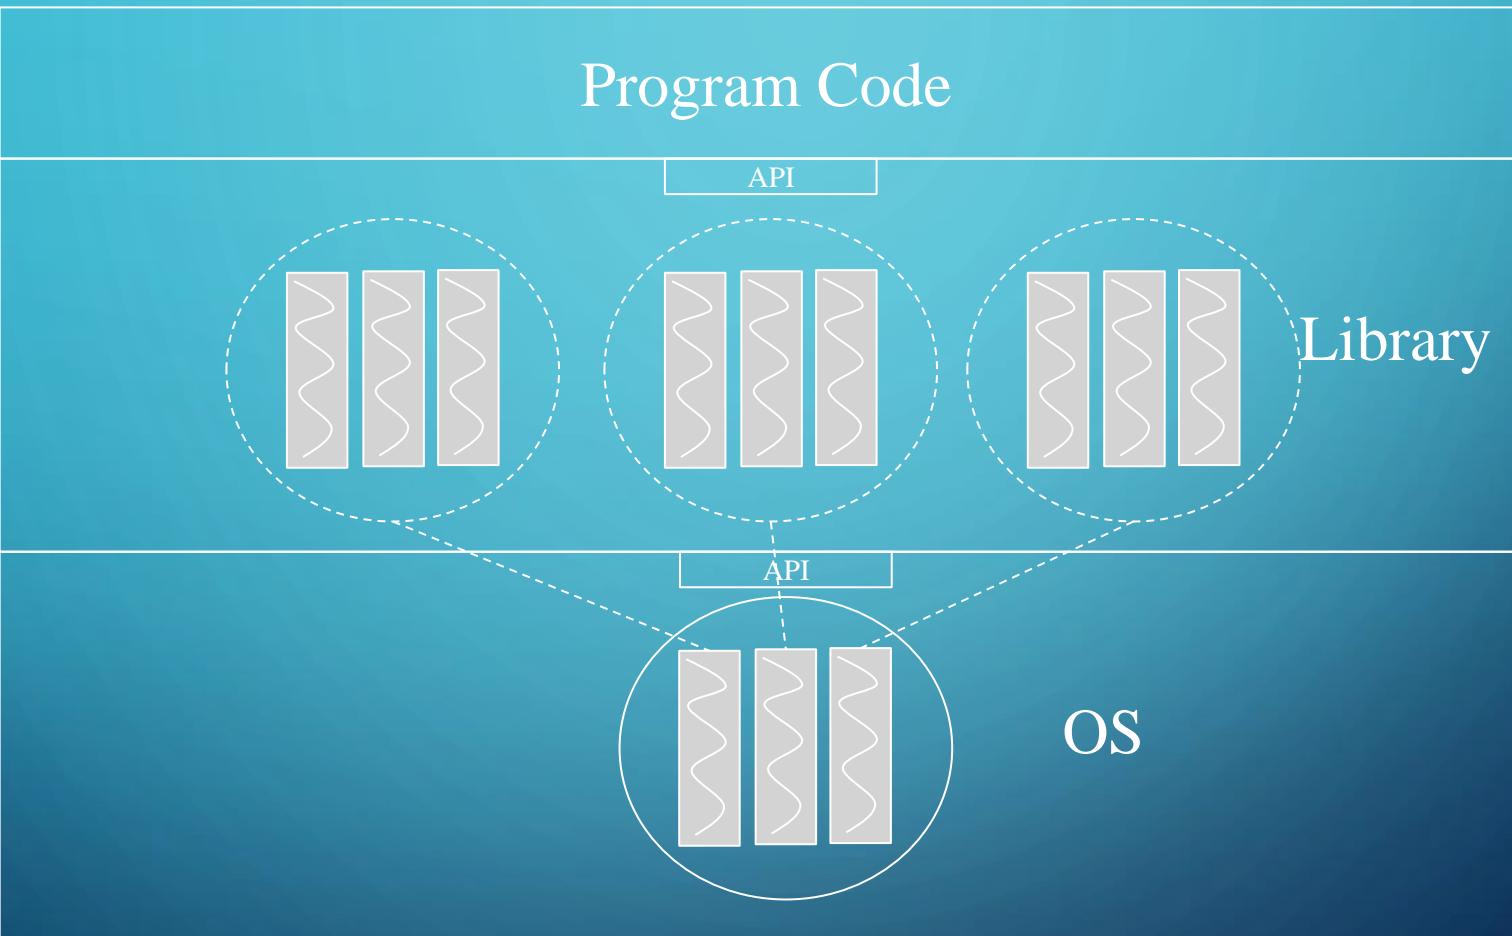
User Level Implementation



ULT and KLT combination



ULT and KLT combination



ULT and KLT comparison

	Advantages	Disadvantages
ULT	<ul style="list-style-type: none">More policies availableManagement operations are more efficient (creation, switching, blocking)	<ul style="list-style-type: none">Program code tends to be more complexThere is no real parallelism
KLT	<ul style="list-style-type: none">Program code is simplerReal parallelism is possible	<ul style="list-style-type: none">Management operations are less efficient

1 / 10 / 100 relationship

- ULT Thread creation: 1 time unit*
- KLT Thread creation: 10 time units*
- Process creation: 100 time units*

CONTENTS

1. Process Management basics
2. Elements of a Process
3. Process lifecycle
4. Threads
- 5. Process and thread scheduling**
6. System services for process and thread management

5- PROCESS AND THREAD SCHEDULING

Scheduling is....

... *to share processor time between processes that can be executed.*

- Scheduling Levels (classically)

- *Short term:* assigns processor to ready processes (5 state model)
- *Medium term:* includes process suspension in secondary memory (7 states model)
- *Long term:* decides which processes are to be admitted to the ready queue

- **Scheduler:** module that decides which process to move from ready to running. (Short-term scheduler)

- **Dispatcher:** module that puts into execution the scheduled process.

SEVEN STATE MODEL

5- PROCESS AND THREAD SCHEDULING



Scheduling in the...
Short-term?
Medium-term?
Long-term?



SEVEN STATE MODEL

5- PROCESS AND THREAD SCHEDULING

- **Short-term scheduler:** decides which of the ready to run processes is going to use the CPU.
- **Medium-term scheduler:** decides which of the processes in memory (ready to run or blocked) is/are going to be swapped out. Its goal is to reduce the multiprogramming degree of the system (the remaining *in memory* processes have available more CPU time and more memory).
- **Long-term scheduler:** decides which of the processes created or swapped out is/are admitted into the memory of the system. It increments the multiprogramming degree of the system.

Metrics about scheduling algorithms performance

For processes (or threads)

- Execution (or return) time (T_e)
- Waiting time (T_w)
- Response time (T_a)

For the System

- Processor use ($C \%$)
- Completed works rate (or Throughput) (P)

General objective

- Minimize process metrics
- Maximize system metrics

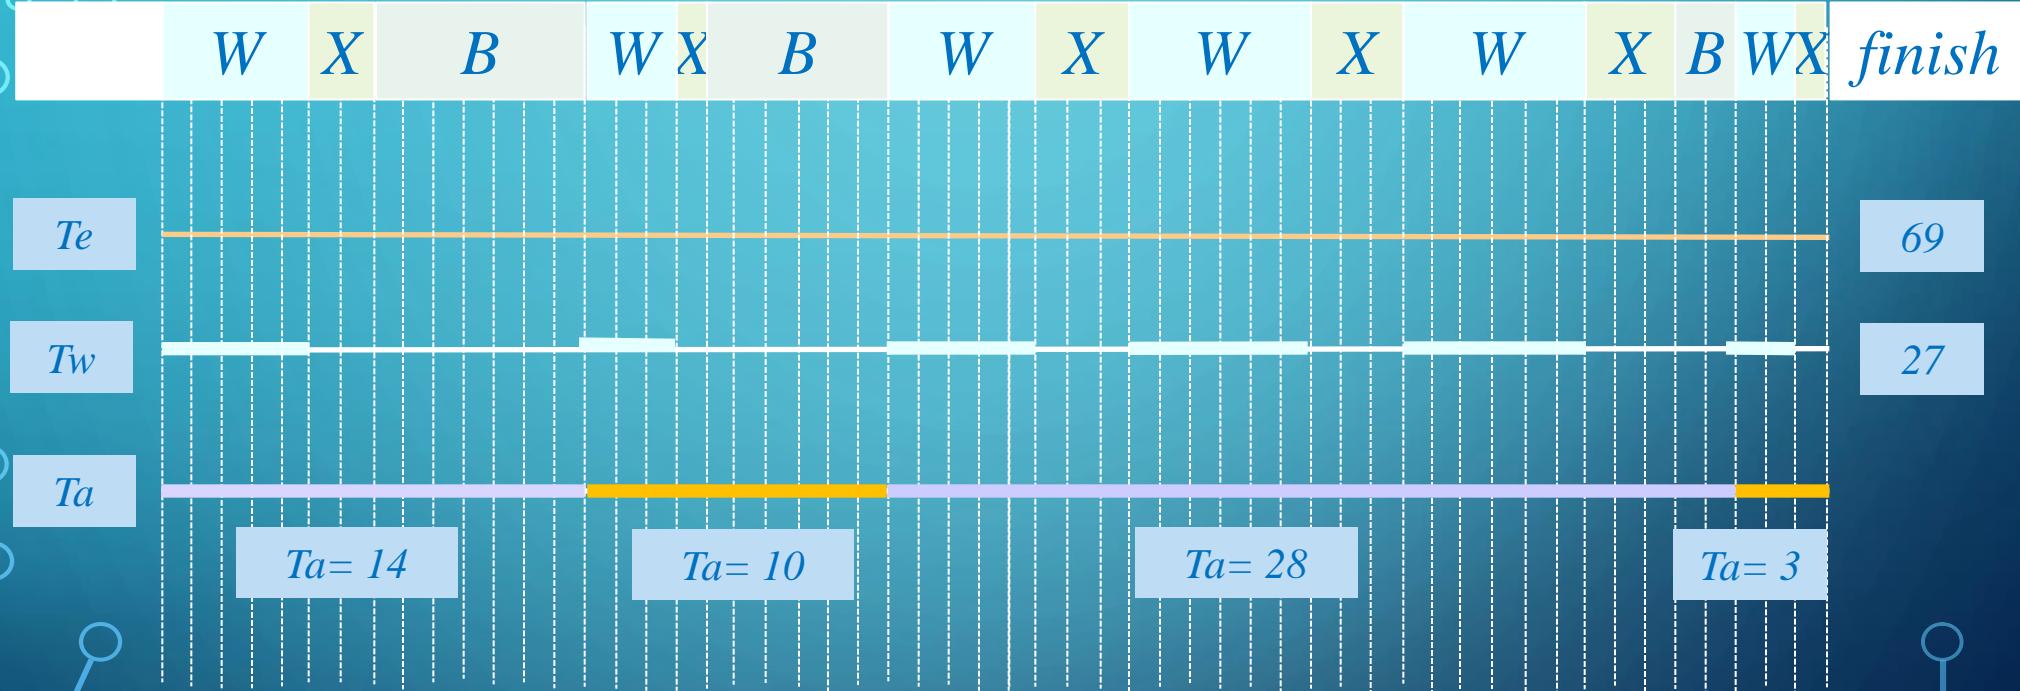
Another objectives...

- Impartiality
- Equitable processor sharing
- Processor efficiency (usage maximization)
- Execution predictability (real time)
- Resource usage balance
- Minimize response time variance
- Reduce context switching time

5.- Process and thread scheduling

Scheduling objectives

P1 lifecycle



Te: Execution Time. From New to Finish.

Tw Waiting Time. Time expended in Ready queue.

Ta Response Time. From a request until the process responds.

5- PROCESS AND THREAD SCHEDULING SCHEDULING POLICIES

- **Non-expulsive**

- FCFS / FIFO (first come, first served) (first in, first out)
- SJF (the shortest, the first)

- **Expulsive / preemptive**

- SRTF (shortest remaining time first)
- Round Robin (rotation shift)

- **+ Priorities**

- Static
- Dynamic
 - With aging
 - Encouraging interactive processes

Current policies: explosive policy mix

Multilevel Queues with dynamic priorities + round robin

*Advantages
disadvantages?*

*Which one would you
choose for your design?*

Why?



FCFS scheduling policy or non-expulsive FIFO

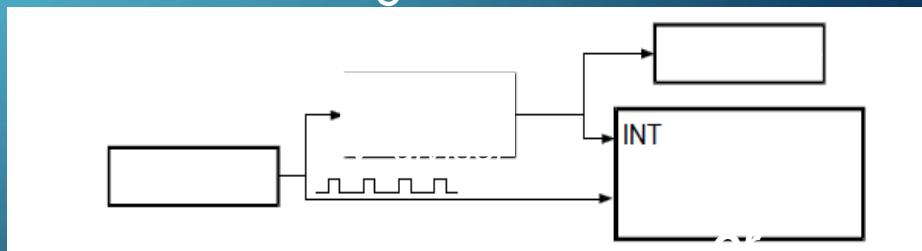
- Ready for execution processes are organized in a FIFO queue
 - When a process changes its state to Ready it goes to the end of the queue.
- A process leaves the CPU when:
 - The process *makes a system call* and as a result it changes to *Blocked*.
 - The process ends.
 - The process *voluntarily leaves* the CPU (*yield*).
- When the CPU is free the scheduler chooses the first process in the ready queue (FIFO).

SJF scheduling policy (shortest first)

- We choose the process with **shorter total duration**.
 - Favors short jobs rather than longer ones.
- Requires knowing the runtime before running the job.
 - *Only useful in production environments where the same jobs are frequently run.*
- Non-expulsive policy (was used in batch processing systems).
- May cause long process starvation.
- Improves the average response time, but increases the variance.
- Deprecated Policy. Designed for batch processing system before multiprogramming.

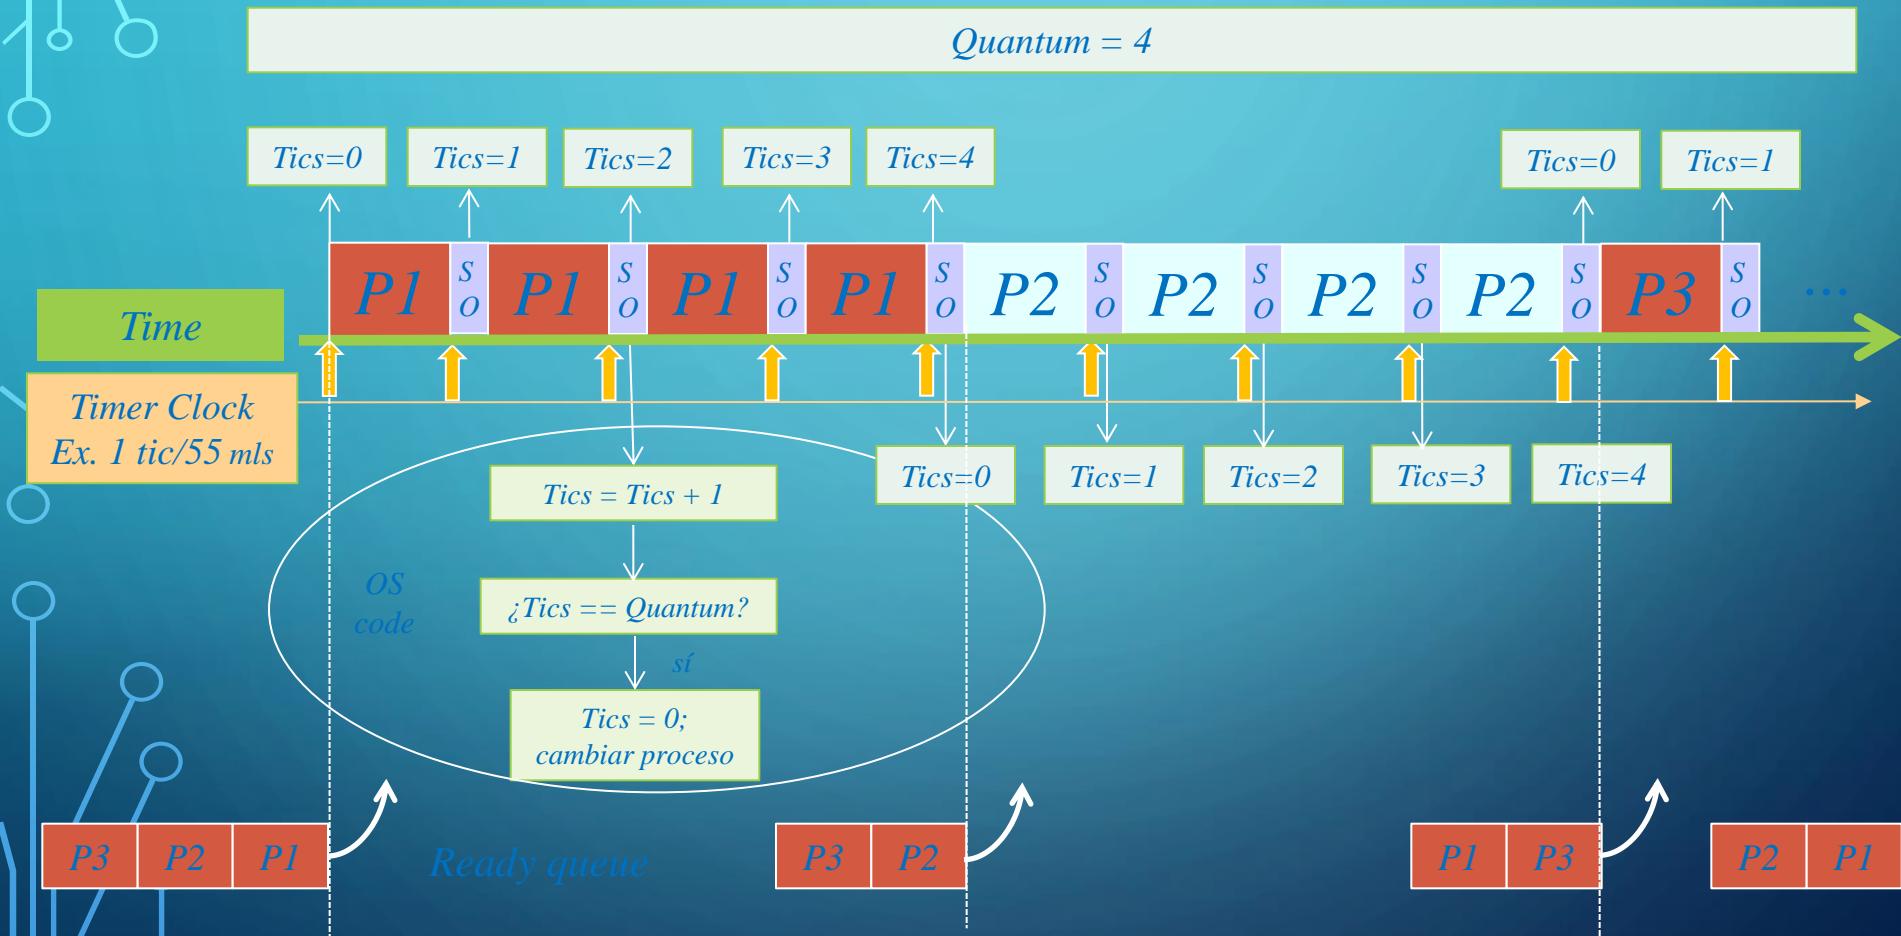
CYCLIC SCHEDULING POLICY OR ROUND ROBIN

- Get a fair allocation of processor time
- The system assigns a *quantum*, or time slice, to the running process
- Ready for execution processes are organized into a FIFO queue
 - When a process changes its state to *Ready* it goes to the queue's tail
- A process leaves the CPU when:
 - The process *makes a system call* and as a result it changes to *Blocked*.
 - The process ends.
 - The process *exhausts its quantum*.
- When the CPU is free, the scheduler chooses *the first process* in the ready queue (FIFO).



Cyclic scheduling policy or Round Robin

Preemption at the end of the quantum



Prioritized scheduling policy

- Each process is assigned a priority:
 - Automatically assigned by the system.
 - Selected by the user.
- The *ready queue* is sorted by priority. The scheduler chooses the highest priority process.
- Types of priority:
 - **Static**
 - **Dynamic**

Prioritized scheduling policy

- **Static priority:**

- **Static:** fixed during the lifetime of the process.
- Easy to implement.
- Introduces little overhead on the system.
- It doesn't adapt to changes in the environment.
- May cause *starvation*.

Prioritized scheduling policy

- Dynamic priority: varies during the process lifetime.
 - More difficult to implement.
 - Enters more overhead on the system.
 - Avoids starvation.
- Example: priority aging.
 - An initial priority is assigned and increases as time passes without running. When running, it returns to the initial priority.
- Example:
 - It increases the priority when the process performs I/O operations and it decreases it when the quantum is exhausted.

Scheduling policy for real-time systems

- The processes have to be executed in predetermined time.
- There are two types of processes:
 - Fixed term: run once at a specific time.
 - Periodic: repeatedly run after a period.
- The moment to run is associated with each process.
- Processes that have not reached their execution time are in a queue, the ones that have reached their execution time go to the ready queue.
- In critical systems each process is assigned a run time slot. The process must be executed before this time limit.

POSIX scheduling policy

- POSIX Scheduling
- Applies to Processes and Threads
- Priority level between 0 and 31
- **It chooses the highest priority thread**
- Available policies:
 - FIFO
 - Round Robin
 - Other

Portable Operating System Interface based on UNIX

5.- Process and thread scheduling

POSIX scheduling policy

- POSIX FIFO Scheduling
 - A FIFO queue per priority.
 - It expels the process if another one with higher priority arrives or if there is an I/O operation.
 - If other (process or thread) with higher priority changes to *Ready state* the current running process is expelled and it becomes the first of its associated priority queue.
 - A *Blocked* process changed to *Ready* goes to the tail of its priority queue.

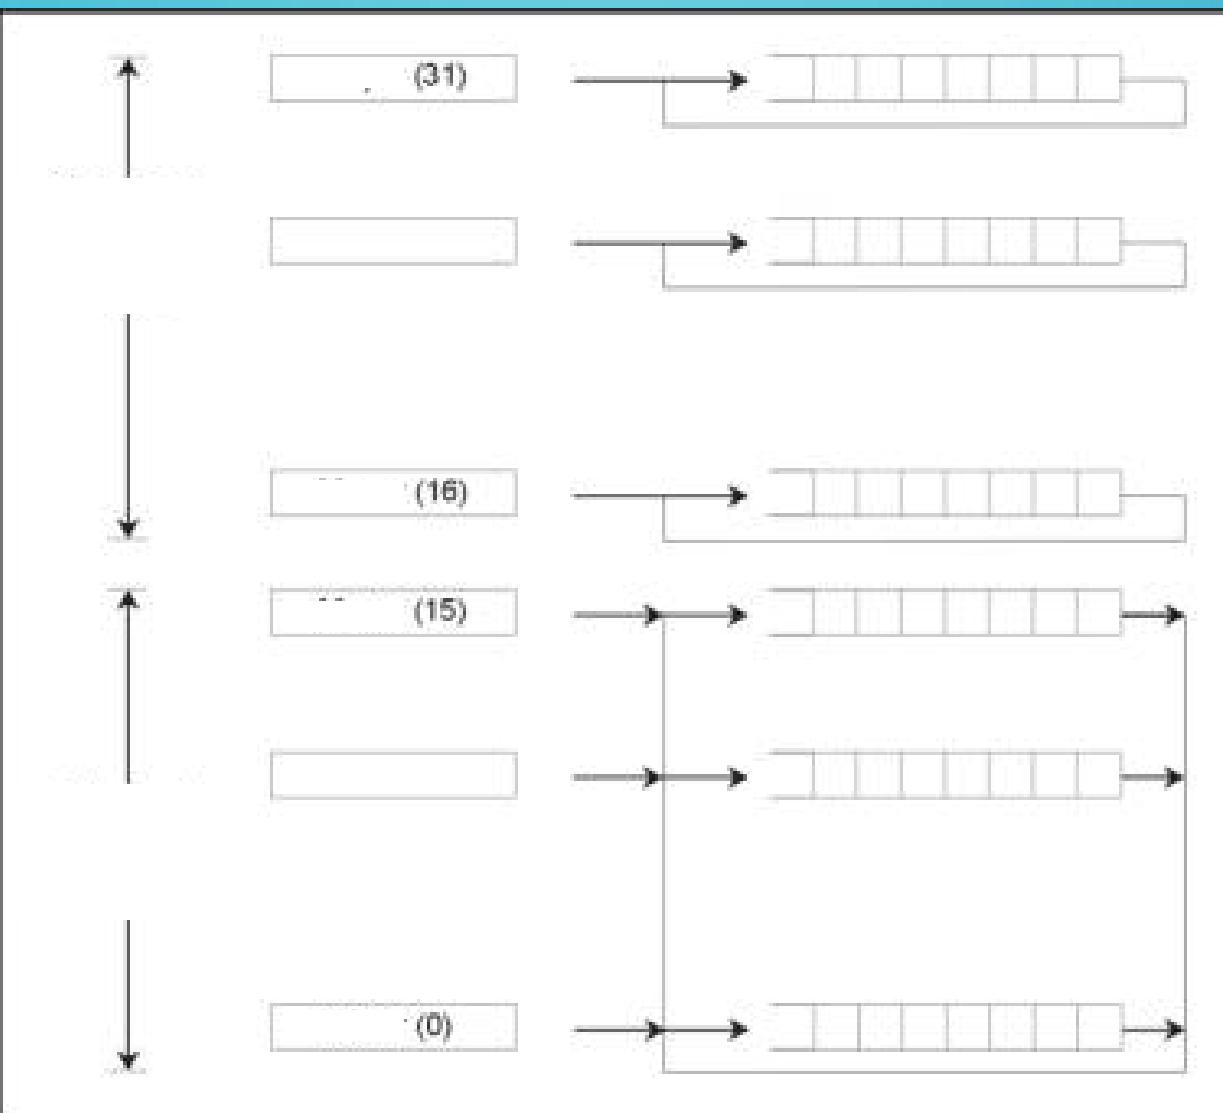
POSIX scheduling policy

- Cyclic Scheduling in POSIX
 - A FIFO queue per priority, processes are assigned a time slice
 - When a process finishes its slice (quantum), it is inserted at the end of its queue
 - When a process is expelled by a higher priority one, it is introduced at the beginning of its queue *without initializing its time slice.*
- Another policy
 - POSIX standard allows to implement a different policy.
- All at once
 - **Each process chooses the policy**, so we have all policies at once.

WINDOWS NT scheduling policy

- Cyclic scheduling with priorities and expulsion.
- Thread scheduling
- 32 scheduling levels (from 0 to 31 maximum)
 - 16 levels with real-time priorities (16-31). Fixed priority, FIFO.
 - 15 levels with varying priorities (1-15). They vary according to the thread behavior. Round Robin.
 - Decremented if the thread runs out of time slice
 - Increased by performing an I/O operation
 - 0 level for system
 - Process and Threads have a base priority

WINDOWS NT scheduling policy



Multiprocessor scheduling

Running the OS on multiprocessors

- **Asymmetric multiprocessing**

- A processor dedicated to running the OS
- It can be slow but the OS is simpler
- OS operations are not concurrent

- **Symmetric Multiprocessing**

- The OS runs on all processors as needed
- Concurrency issues inside the OS services
- OS code is more complex
- More efficient execution

Multiprocessor scheduling

Scheduling criteria

- **Processor affinity**

- Tendency to assign the same processor to a process for maximizing the cache information
 - (each processor has a cache)
- Strict affinity
 - Possibility of indicating on the affinity process to a processor or group (for critical processes)

- **Load sharing**

- Load Balancing processor execution.

Multiprocessor scheduling

Two possibilities for process scheduling

- Unique queue (for Ready processes) for all processors
- One queue per processor

Multiprocessor scheduling

Single queue scheduling

- If a processor is free:
 - Select a process from the queue
 - By priority
 - By affinity (is affine if that processor has executed it the last time)
- If there is NOT free processor:
 - A process changed to *Ready* (at birth or from *Blocked*) is assigned to a processor IF an affine processor executes a lower priority process (that is expelled)

Advantage → Load Balancing

Drawback → Single queue needs protected access, implies bottleneck

Multiprocessor scheduling

One queue per processor scheduling

- Each process is assigned to a processor
- Each processor executes process from its own queue

Advantage

- Single queue bottleneck disappears

Drawbacks

- Load Rebalancing is not automatic
 - Periodically
 - If a processor is freed
- Processes are changed from one queue to another.
- It takes into account affinities (strict and cache)

Multiprocessor scheduling

Thread scheduling policies

- Not taken into account the relationships between same process' threads, OR
- It takes into account the relationships between threads
 - Shared time
 - Applications are multiplexed in time
 - Time is assigned to each application (set of threads) in a processor. All threads are running on the same processor.
 - Shared space (co-scheduling)
 - Applications are multiplexed in space and time
 - Application's threads are distributed among groups of processors
 - Applications are also multiplexed in the time

Both types can use shared or local queues

Multiprocessor scheduling

Current Systems

- Symmetric Multiprocessing
 - The OS runs on all processors
- Scheduling
 - Windows (client) → One global queue
 - Windows (server) → Local queues
 - Linux → Local queues

CONTENTS

1. Process Management basics
2. Elements of a Process
3. Process lifecycle
4. Threads (or lightweight Processes)
5. Process and thread scheduling
6. **System services for process and thread management**

6. SYSTEM SERVICES FOR PROCESS AND THREAD MANAGEMENT

POSIX: FORK() SYSTEM CALL

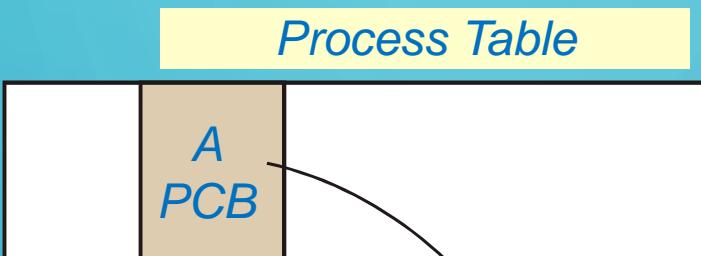
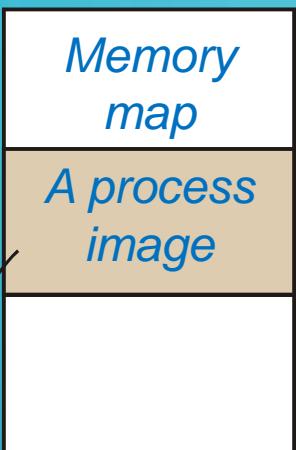
Services for Process Management

- Process Creation
 - **fork, execl, execv, execle, execve, execlp, execvp**
- Process Termination
 - **exit, wait, waitpid**
- Process identification
 - **getpid, getppid, getuid, geteuid, getgid, getegid,**
- The environment of a process
 - **getenv**

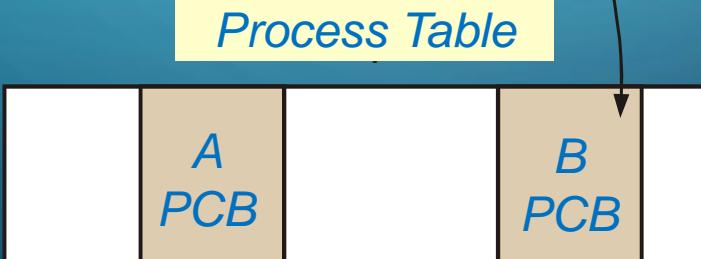
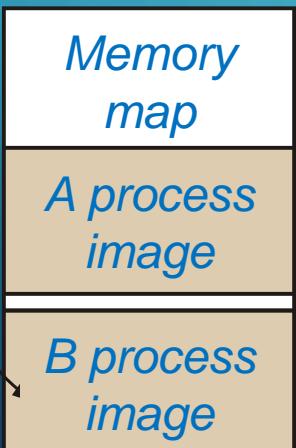
Portable Operating System Interface based on UNIX

6. SYSTEM SERVICES FOR PROCESS AND THREAD MANAGEMENT

POSIX: FORK() SYSTEM CALL



'A' process performs a `fork()` system call
which creates a child process 'B'



New PID
New process image
Different return value (0 in child)

6. SYSTEM SERVICES FOR PROCESS AND THREAD MANAGEMENT

POSIX: FORK() SYSTEM CALL

Syntax: `pid_t fork();`

Operation: Clones the process making the call

Implications:

- The process making the call becomes the parent. The new process is the child.
- Process image and the PCB are duplicated
- Parent and child continue execution at the same point
- Some PCB's child values will be different
 - PID of the child
 - Addresses of the memory segments of the process image
 - Execution time counters and signals in the child are reset
- `fork()` return value is *0 in the child and the child's PID in the parent*

6. SYSTEM SERVICES FOR PROCESS AND THREAD MANAGEMENT

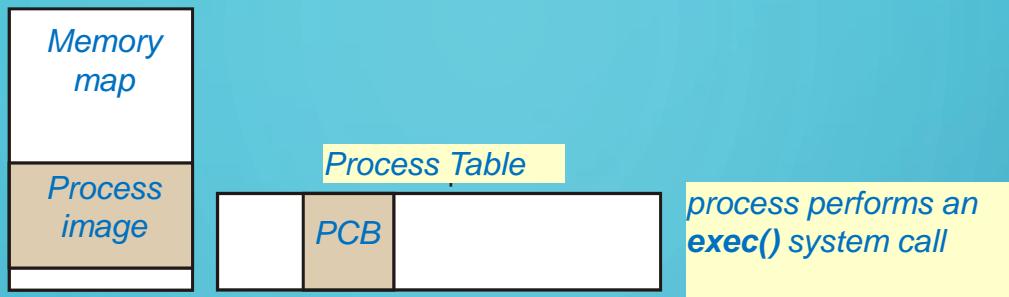
POSIX: FORK() SYSTEM CALL

Example

```
main ()  
{  
    if (fork () != 0)  
        printf ("I am the parent\n");  
    else  
        printf ("I am the child\n");  
}
```

6. SYSTEM SERVICES FOR PROCESS AND THREAD MANAGEMENT

POSIX: FORK() SYSTEM CALL



POSIX : EXEC() SYSTEM CALL

Syntax: exec()

Operation: Change the program that is running a process by another one indicated as parameter (an executable file).

Implications:

- Liberation of the process image memory
- Reassigning new address space to load the executable file
- Fill the PCB with initial values for registers and new memory segments

6. SYSTEM SERVICES FOR PROCESS AND THREAD MANAGEMENT

POSIX : EXEC() SYSTEM CALL

Example

```
main()
{
    if (fork() == 0)
        execlp("ls", "ls", "-l", NULL);
    printf("The parent continues alive\n");
}
```

Command interpreter (Shell)

Repeat

read command line

do fork()

If this process is the child do exec(command line)

If this process is the parent

If the command line **does not** end with "&"

 wait() until child finishes

Until command line = "exit"

EXERCISES

```
void main() {  
    int a = 0;  
    if (fork() != 0) {  
        a = a+1;  
        printf(a);  
    }  
}
```

```
void main() {  
    int a = 0;  
    if (fork() != 0) {  
        a = a+1;  
    }  
    printf(a);  
}
```

What is the result of its execution?



EXERCISES

How many processes are created when you run this code?

```
for (i=0; i < n; i++) {  
    fork();  
    execvp("ls", "ls", "-l", NULL)  
}
```

```
for (i=0; i < n; i++) {  
    fork();  
}  
execvp("ls", "ls", "-l", NULL)
```



EXERCISES

How many processes are created when you run this code?

```
for (i=0; i < n; i++) {
    if (! fork()) execvp("ls", "ls", "-l", NULL)
}
```

```
for (i=0; i < n; i++) {
    if (fork()) execvp("ls", "ls", "-l", NULL)
}
```



EXERCISES

Which of the following statements are correct?

- a) The content of a memory image of a process remains unchanged throughout the execution
- b) The process execution stack is used to store local variables and parameters each time a method (or function) is invoked
- c) The program code that has been loaded into memory is part of the memory image of the process
- d) None of the above

EXERCISES

According to the basic model lifecycle (New-Ready-Running-Blocked-Finished)

Which of the following statements are correct?

- a) A running process does not always have a PCB
- b) A running process must go through every state of its lifecycle before it finishes
- c) The Ready processes queue stores the PCBs of the Ready state processes
- d) None of the above

EXERCISES

According to the basic model lifecycle (New-Ready-Running-Blocked-Finished)

Which of the following statements are correct?

- a) A running process will change to Ready state when executing a system call that involves a slow I/O operation
- b) When a process is born, the system must assign a PCB, which stores the process image, among other things
- c) A Blocked process will change to Running when the event it was waiting for finally takes place
- d) None of the above

EXERCISES

Which of the following statements are correct?

- a) When an interruption occurs, the hardware keeps track of the Program Counter (PC) and places on it the OS routine address for interruption handling
- b) The interruption handling routine keeps the remaining registers before it manages the interruption
- c) In case an interruption triggers a process change, the scheduler and dispatcher will be called
- d) None of the above